



Quantitative techniques in computer design

Nilanjan Byabarta

UEM Kolkata



Agenda

measuring and reporting performance

Performance Measurement

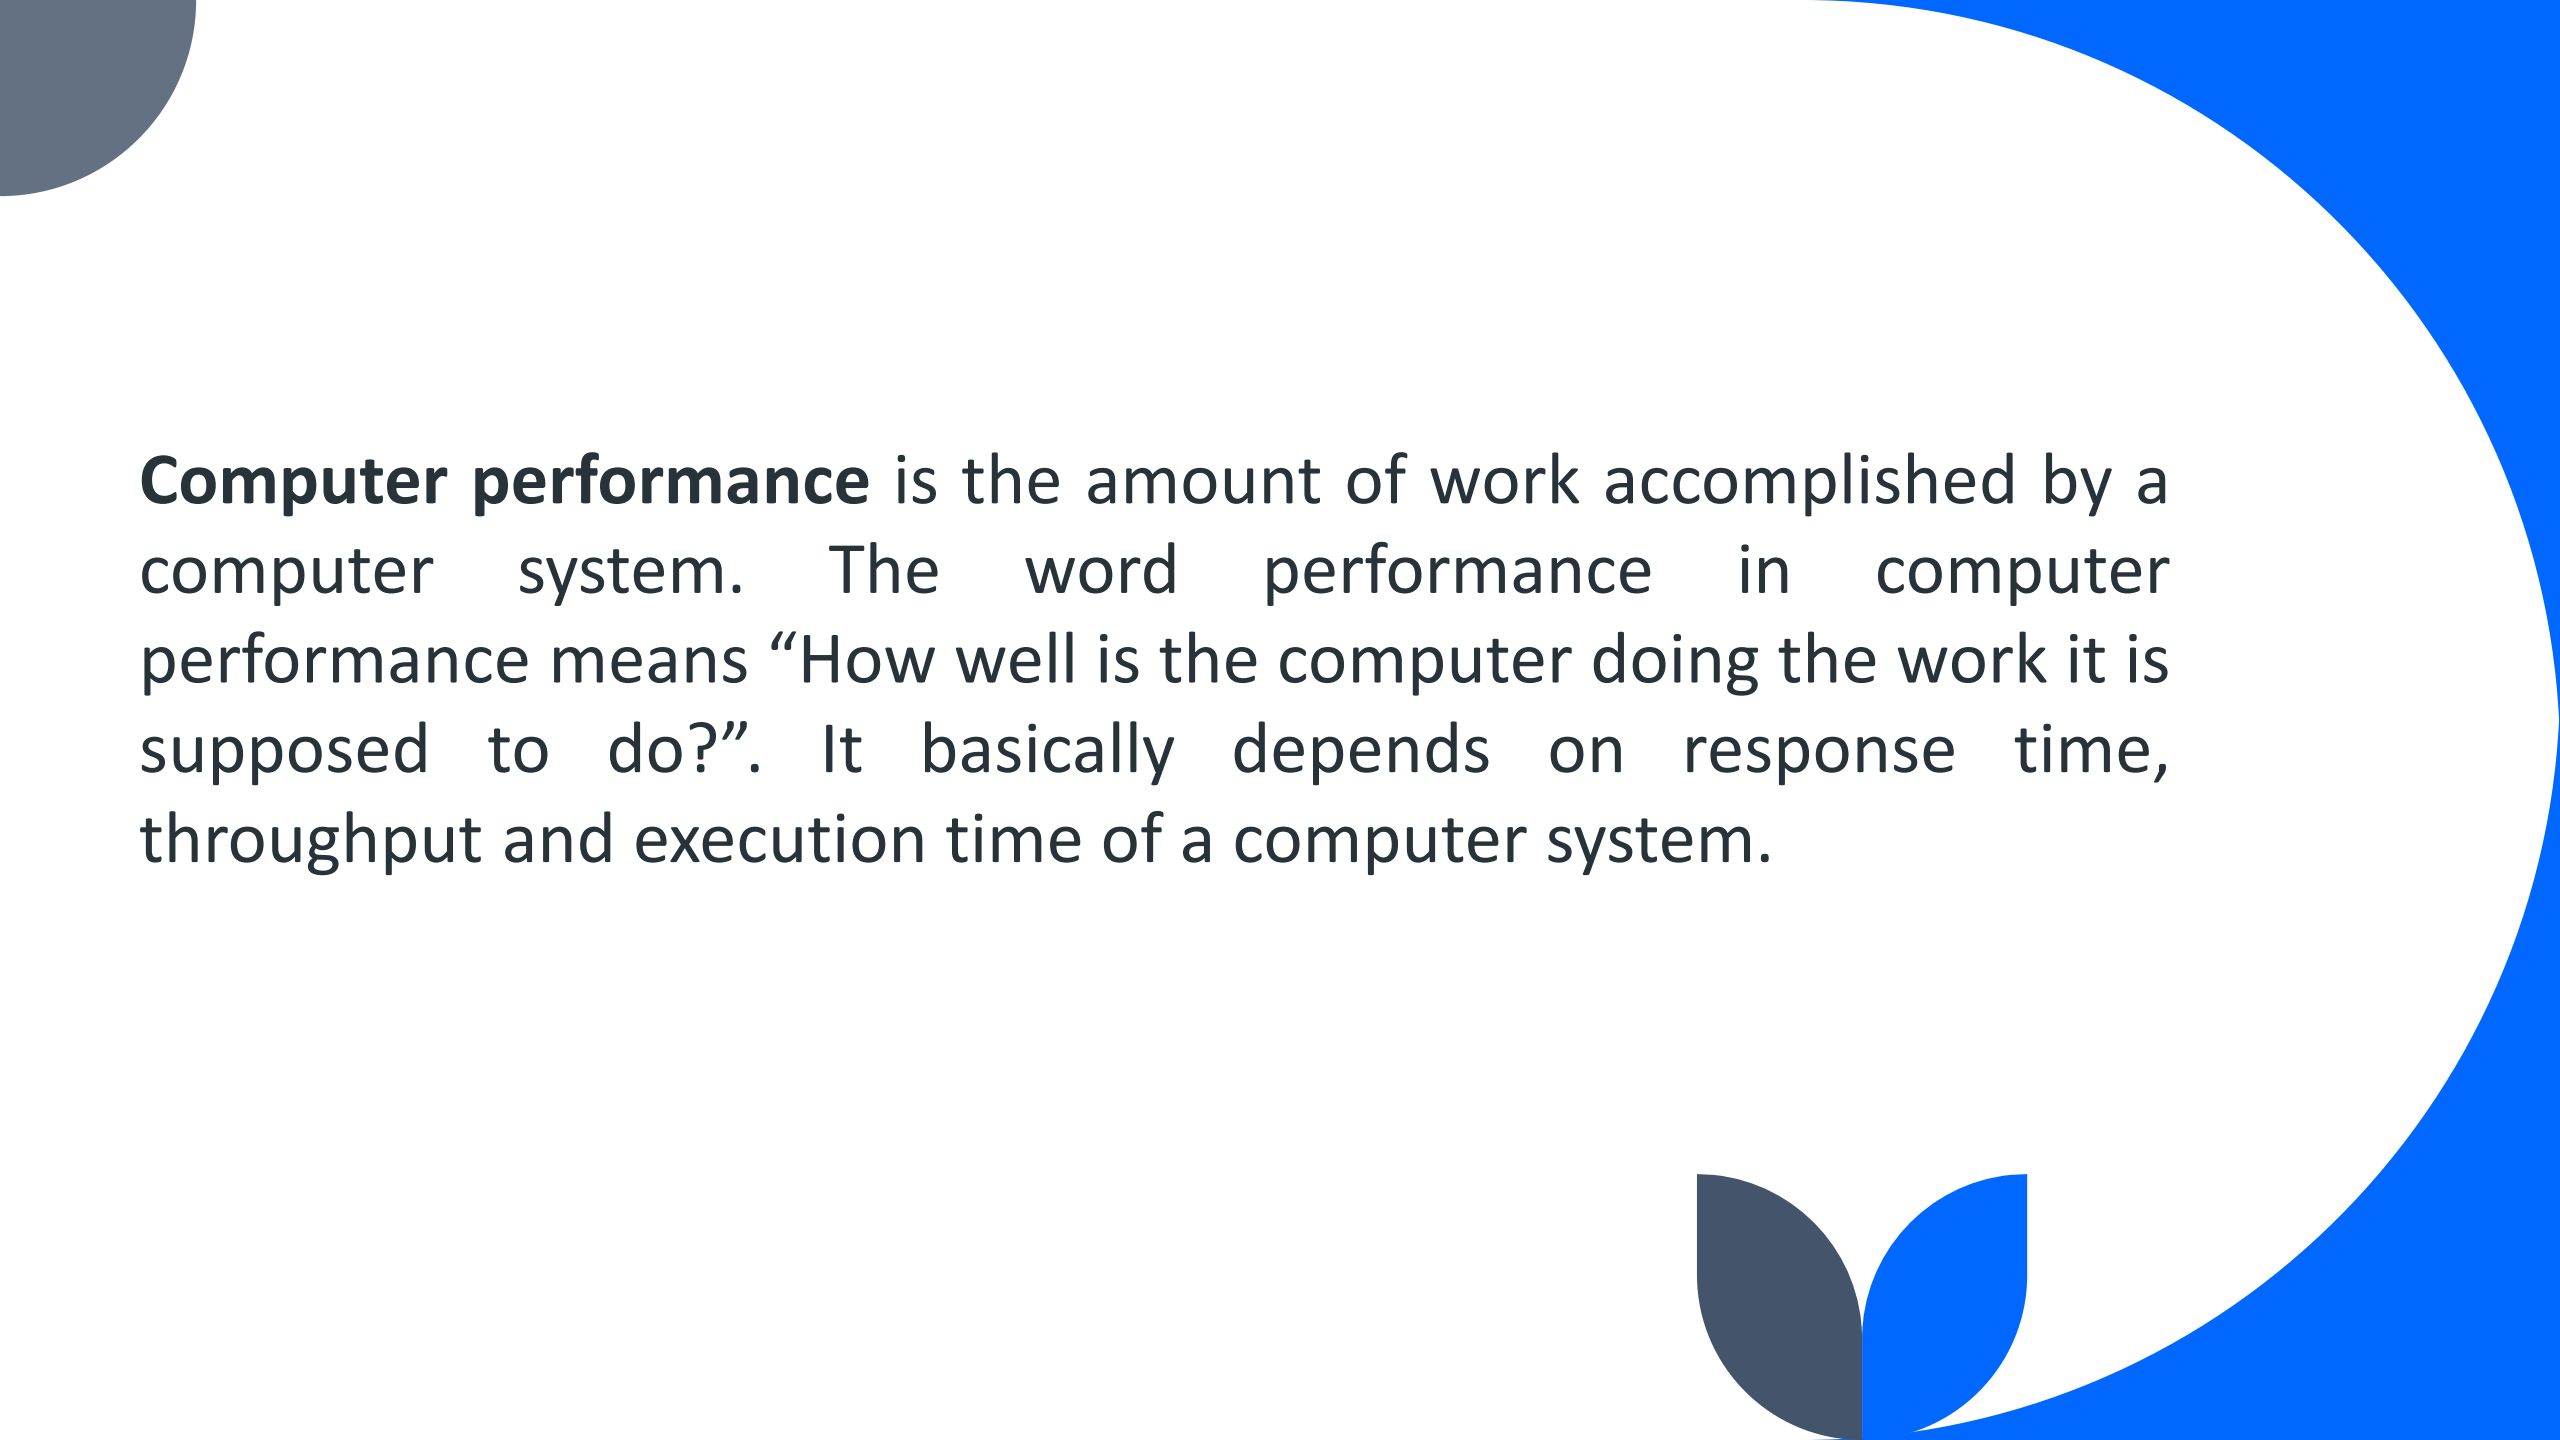
Recall performance is function of

- CPI: cycles per instruction

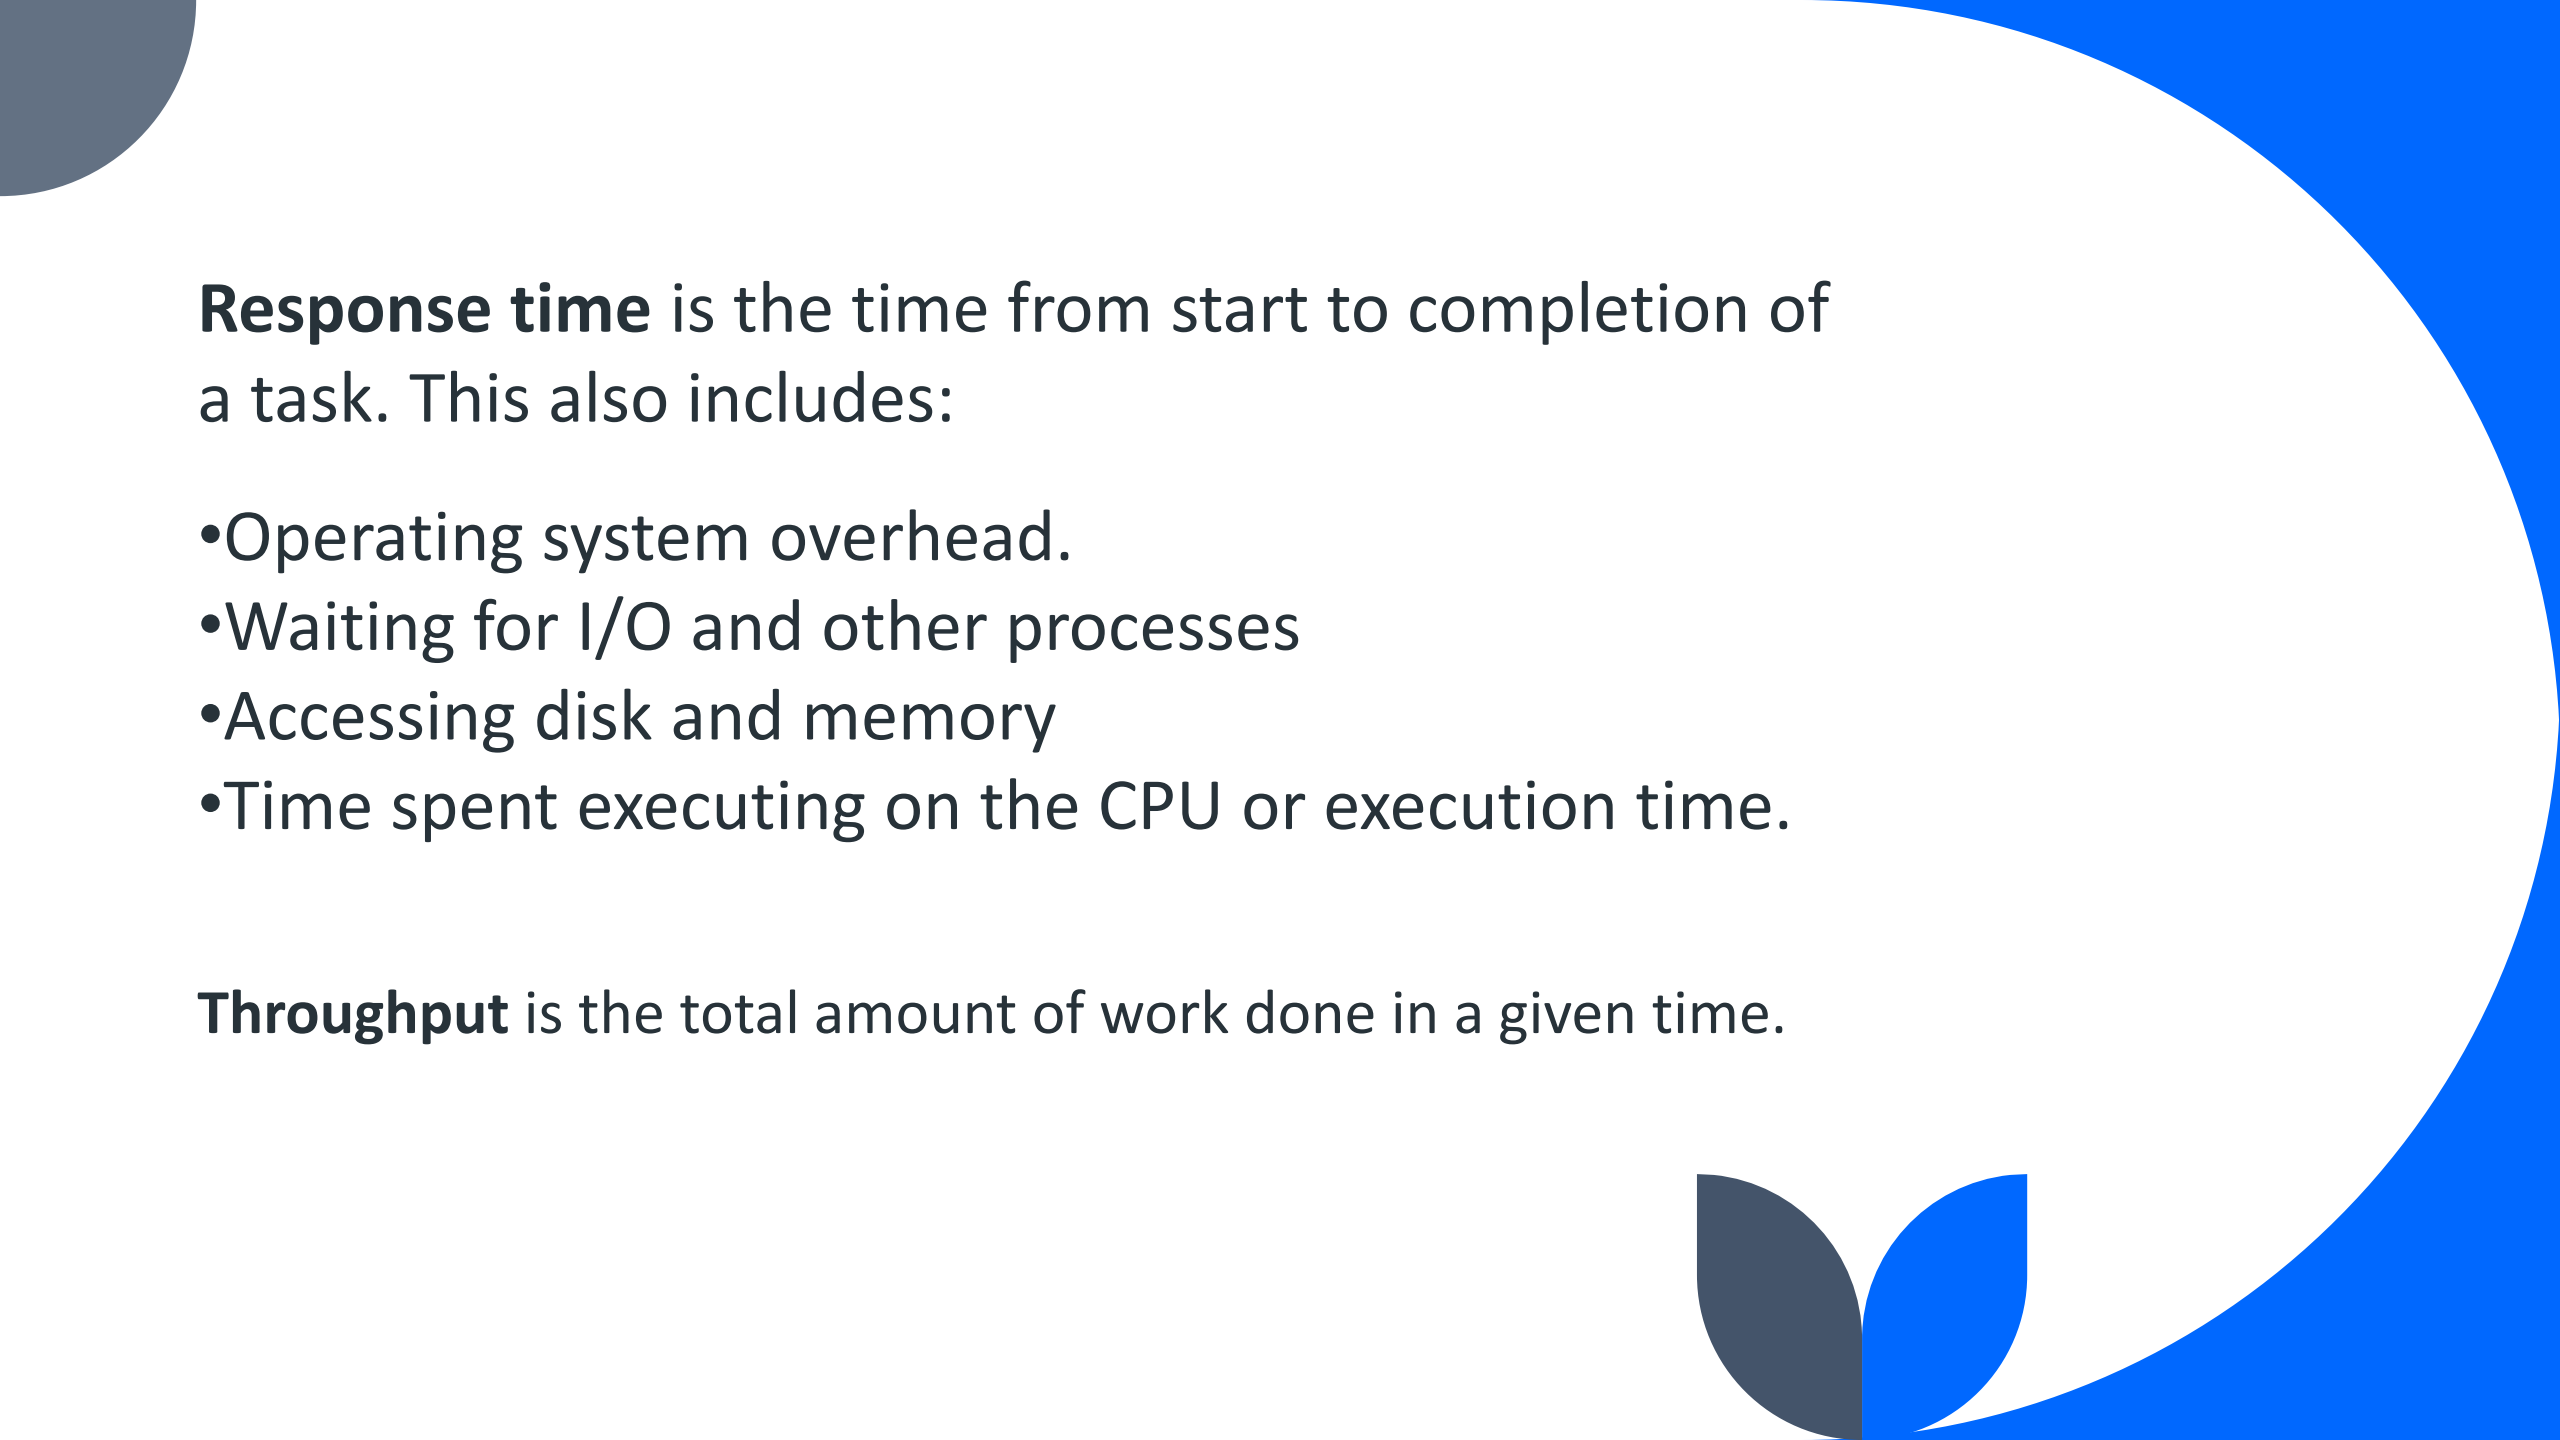
- Clock cycle

- Instruction count

Reducing any of the 3 factors will lead to improved performance



Computer performance is the amount of work accomplished by a computer system. The word performance in computer performance means “How well is the computer doing the work it is supposed to do?”. It basically depends on response time, throughput and execution time of a computer system.



Response time is the time from start to completion of a task. This also includes:

- Operating system overhead.
- Waiting for I/O and other processes
- Accessing disk and memory
- Time spent executing on the CPU or execution time.

Throughput is the total amount of work done in a given time.

CPU execution time is the total time a CPU spends computing on a given task. It also excludes time for I/O or running other programs. This is also referred to as simply CPU time.

Performance is determined by execution time as performance is inversely proportional to execution time.

Performance = $(1 / \text{Execution time})$

And,

$(\text{Performance of A} / \text{Performance of B}) = (\text{Execution Time of B} / \text{Execution Time of A})$

Execution time = CPU clock cycles x clock cycle time

Since clock cycle time and clock rate are reciprocals, so,

Execution time = CPU clock cycles / clock rate

The number of CPU clock cycles can be determined by,

CPU clock cycles = (No. of instructions / Program) x (Clock cycles / Instruction) = Instruction Count x CPI

Execution time = Instruction Count x CPI x clock
cycle time

= Instruction Count x CPI / clock rate

How to Improve Performance?

To improve performance you can either:


- Decrease the CPI (clock cycles per instruction) by using new Hardware.
- Decrease the clock time or Increase clock rate by reducing propagation delays or by use pipelining.
- Decrease the number of required cycles or improve ISA or Compiler.

A Few Words About Where We Are Headed

Performance = 1 / Execution time simplified to 1 / CPU execution time


CPU execution time = Instructions × CPI / (Clock rate)

Performance = Clock rate / (Instructions × CPI)





Try to achieve CPI = 1
with clock that is as
high as that for CPI > 1
designs; is CPI < 1
feasible?

Design memory & I/O
structures to support
ultrahigh-speed CPUs



Define an instruction set;
make it simple enough
to require a small number
of cycles and allow high
clock rate, but not so
simple that we need many
instructions, even for very
simple tasks



Design hardware
for CPI = 1; seek
improvements with
CPI > 1

Design ALU for
arithmetic & logic
ops



Performance Measurement

First step is to apply concept of pipelining to the instruction execution process

- Overlap computations

What does this do?

- Decrease clock cycle

- Decrease effective CPU time compared to original clock cycle

Introduction

Principles that are useful in design and analysis of computers:

- Make the common case fast !
 - If a design trade-off is necessary, favor the frequent case (which is often simpler) over the infrequent case.
 -
 - For example, given that overflow in addition is infrequent, favor optimizing the case when no overflow occurs.

Introduction

Typical performance metrics:

Response time

Throughput

Speedup of X relative to Y

$\text{Execution time}_Y / \text{Execution time}_X$

Execution time

Wall clock time: includes all system overheads

CPU time: only computation time

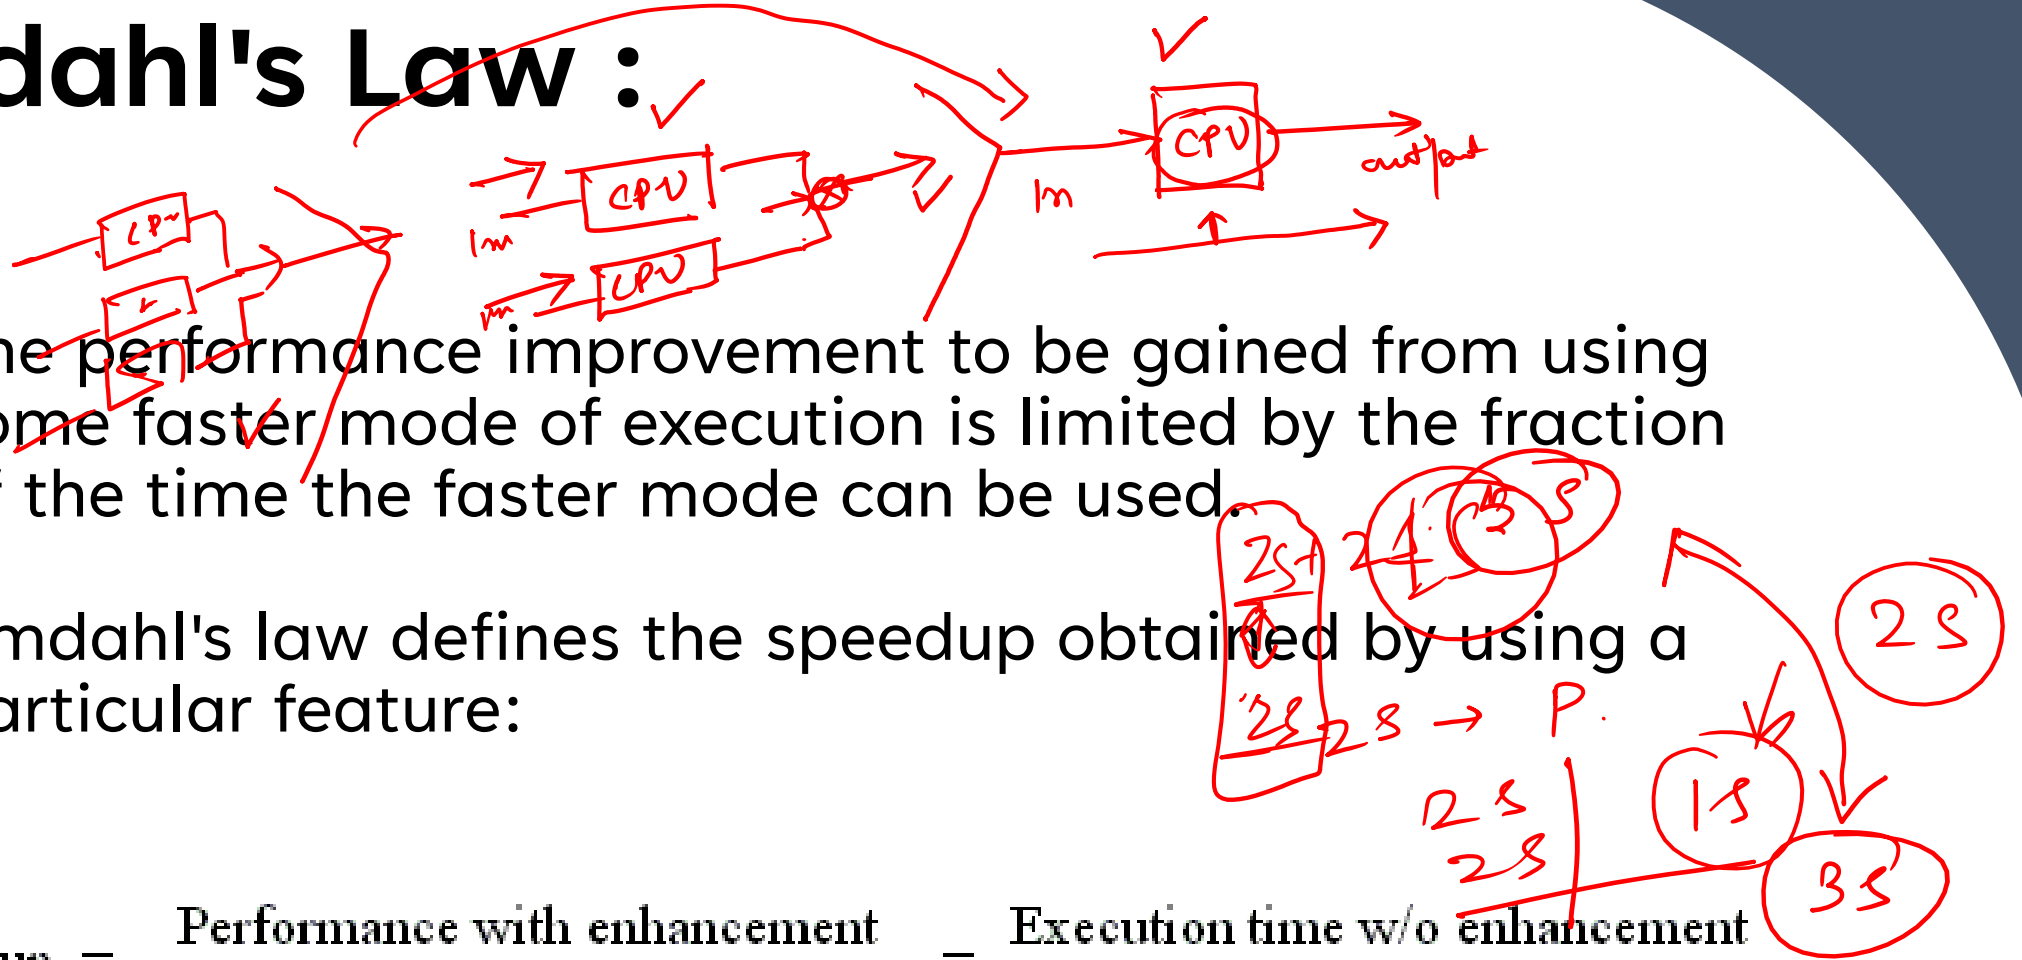
Objective:

- Determine the frequent case.
- Determine how much improvement in performance is possible by making it faster.
- Amdahl's law can be used to quantify the latter given that we have information concerning the former.


Amdahl's Law :

- The performance improvement to be gained from using some faster mode of execution is limited by the fraction of the time the faster mode can be used.
- Amdahl's law defines the speedup obtained by using a particular feature:

$$\text{Speedup} = \frac{\text{Performance with enhancement}}{\text{Performance without enhancement}} = \frac{\text{Execution time w/o enhancement}}{\text{Execution time with enhancement}}$$



Areas of growth



- Two factors:
 - $\text{Fraction}_{\text{enhanced}}$: Fraction of compute time in original machine that can be converted to take advantage of the enhancement.
 - Always ≤ 1 .
 - $\text{Speedup}_{\text{enhanced}}$: Improvement gained by enhanced execution mode:

$$\frac{\text{Time of original mode}}{\text{Time of enhanced mode}}$$

Always > 1

• ”

- Execution time using original machine with enhancement:

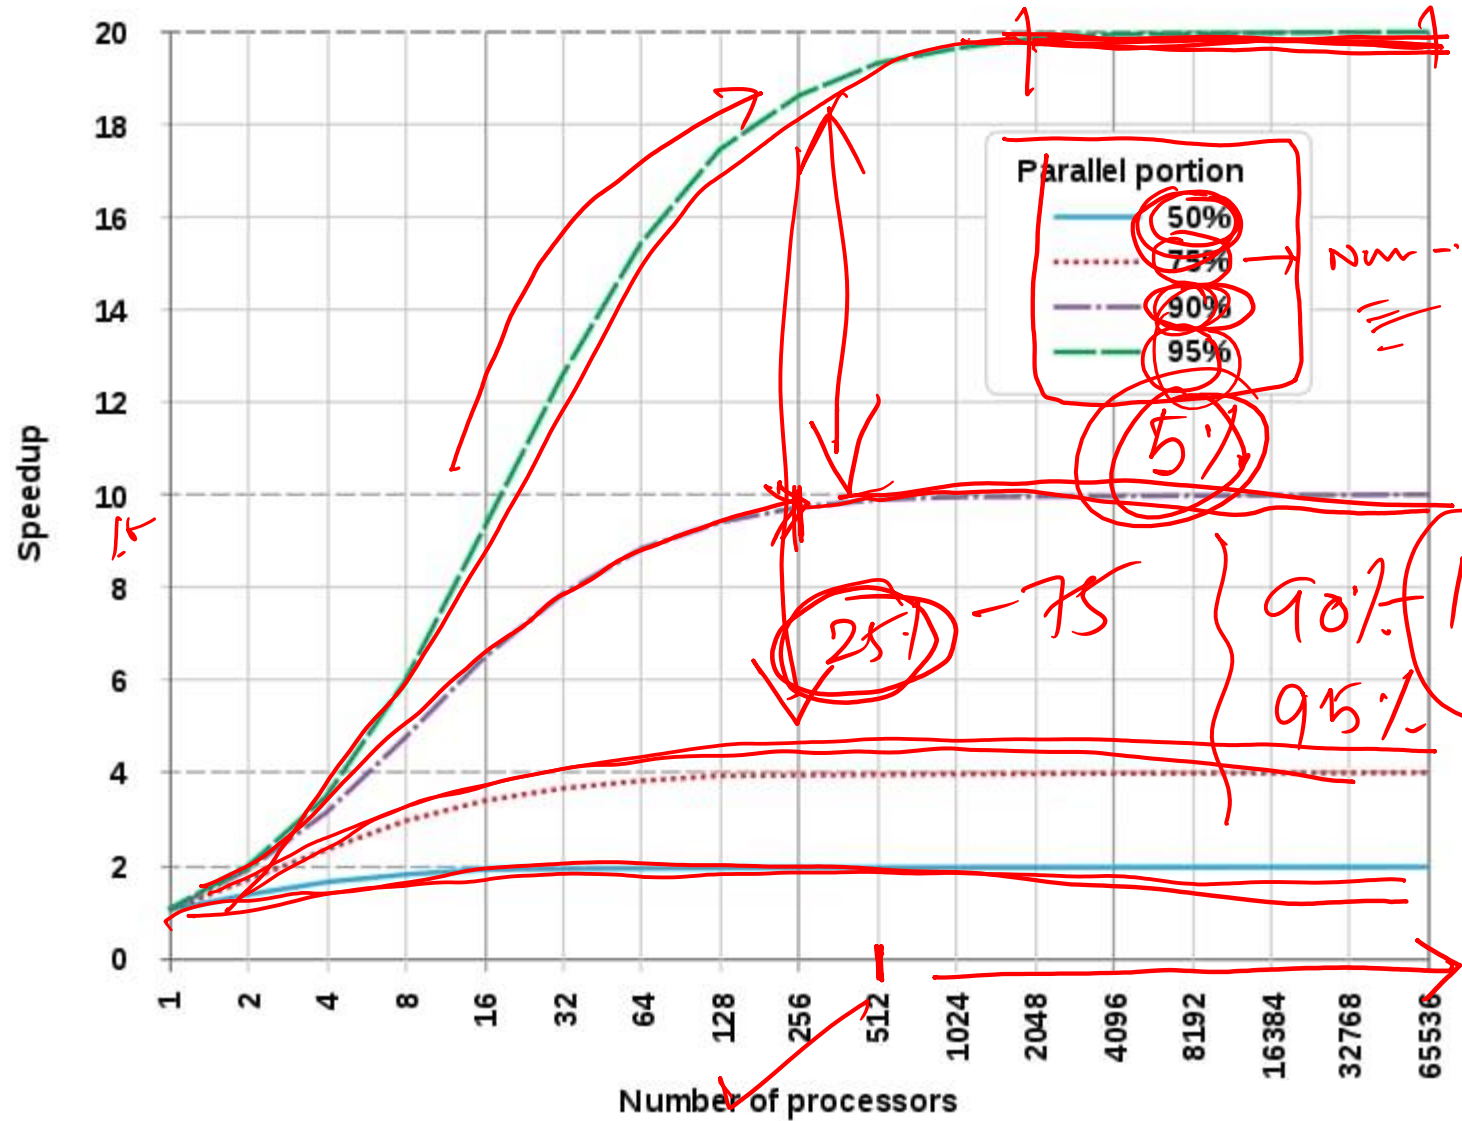
$$\text{Exec time}_{new} = \text{Exec time}_{old} \times \left((1 - \text{Fraction}_{enhanced}) + \frac{\text{Fraction}_{enhanced}}{\text{Speedup}_{enhanced}} \right)$$

Time spent using UNenhanced mode

Time spent using enhancement

$$= T \times \left\{ (1 - P) + \frac{P}{S} \right\}$$

Amdahl's Law



- *Speedup*_{overall} using Amdahl's Law:

$$Speedup_{overall} = \frac{ExecTime_{old}}{ExecTime_{new}} = \frac{1}{\left((1 - Fraction_{enhanced}) + \frac{Fraction_{enhanced}}{Speedup_{enhanced}} \right)}$$

CPU Performance Equation:

- Often it is difficult to measure the improvement in time using a new enhancement directly.
- A second method that decomposes the CPU execution time into three components makes this task simpler.
- CPU Performance Equation:

$$\text{CPU time} = \text{CPU clock cycles for a program} \times \text{Clock cycle time}$$

$$\text{CPU time} = \frac{\text{CPU clock cycles for a program}}{\text{Clock rate}}$$

• *where, for example, Clock cycle time = 2ns for a 500MHz Clock rate.*

CPU Performance Equation:

- An alternative to "number of clock cycles" is "number of instructions executed" or Instruction Count (IC).
- Given both the "number of clock cycles" and IC of a program, the average Clocks Per Instruction (CPI) is given by:

$$\text{CPI} = \frac{\text{CPU clock cycles of a program}}{\text{IC}}$$

$$\text{CPU time} = \text{IC} \times \text{CPI} \times \text{Clock cycle time} = \frac{\text{IC} \times \text{CPI}}{\text{Clock rate}}$$

CPU Performance Equation:

- One difficulty: It is difficult to change one in isolation of the others:
- Clock cycle time: Hardware and Organization.
- CPI: Organization and Instruction set architecture.
- Instruction count: Instruction set architecture and Compiler technology.

$$\text{CPU time} = \left(\sum_{i=1}^n \text{CPI}_i \times \text{IC}_i \right) \times \text{Clock cycle time}$$

•where IC_i represents number of time instruction i is executed in a program and CPI_i represents the average number of clock cycles for instruction i .

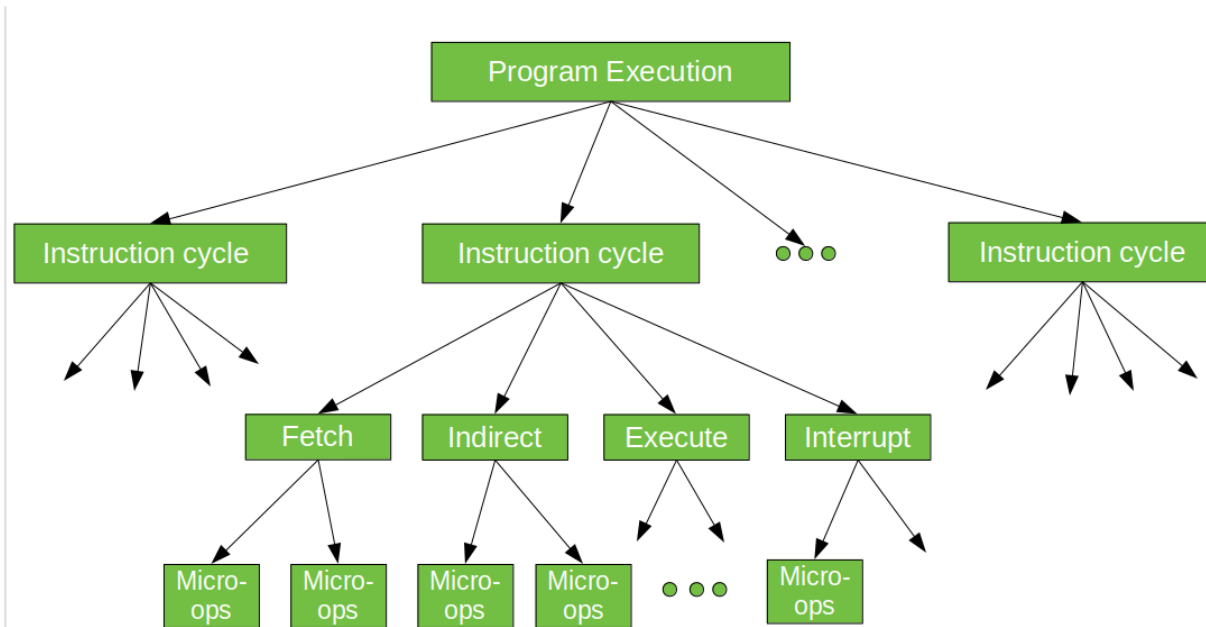
•Why isn't CPI_i a constant

**(million instruction per second)
is NOT an alternative metric to time.**

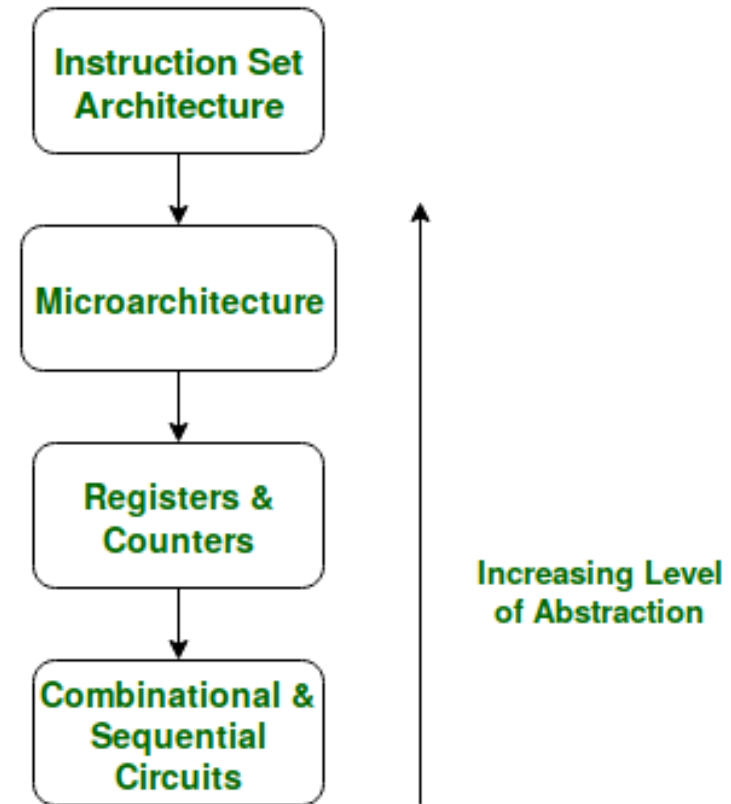
$$\text{MIPS} = \frac{\text{Instruction count}}{\text{Exec Time} \times 10^6} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}$$

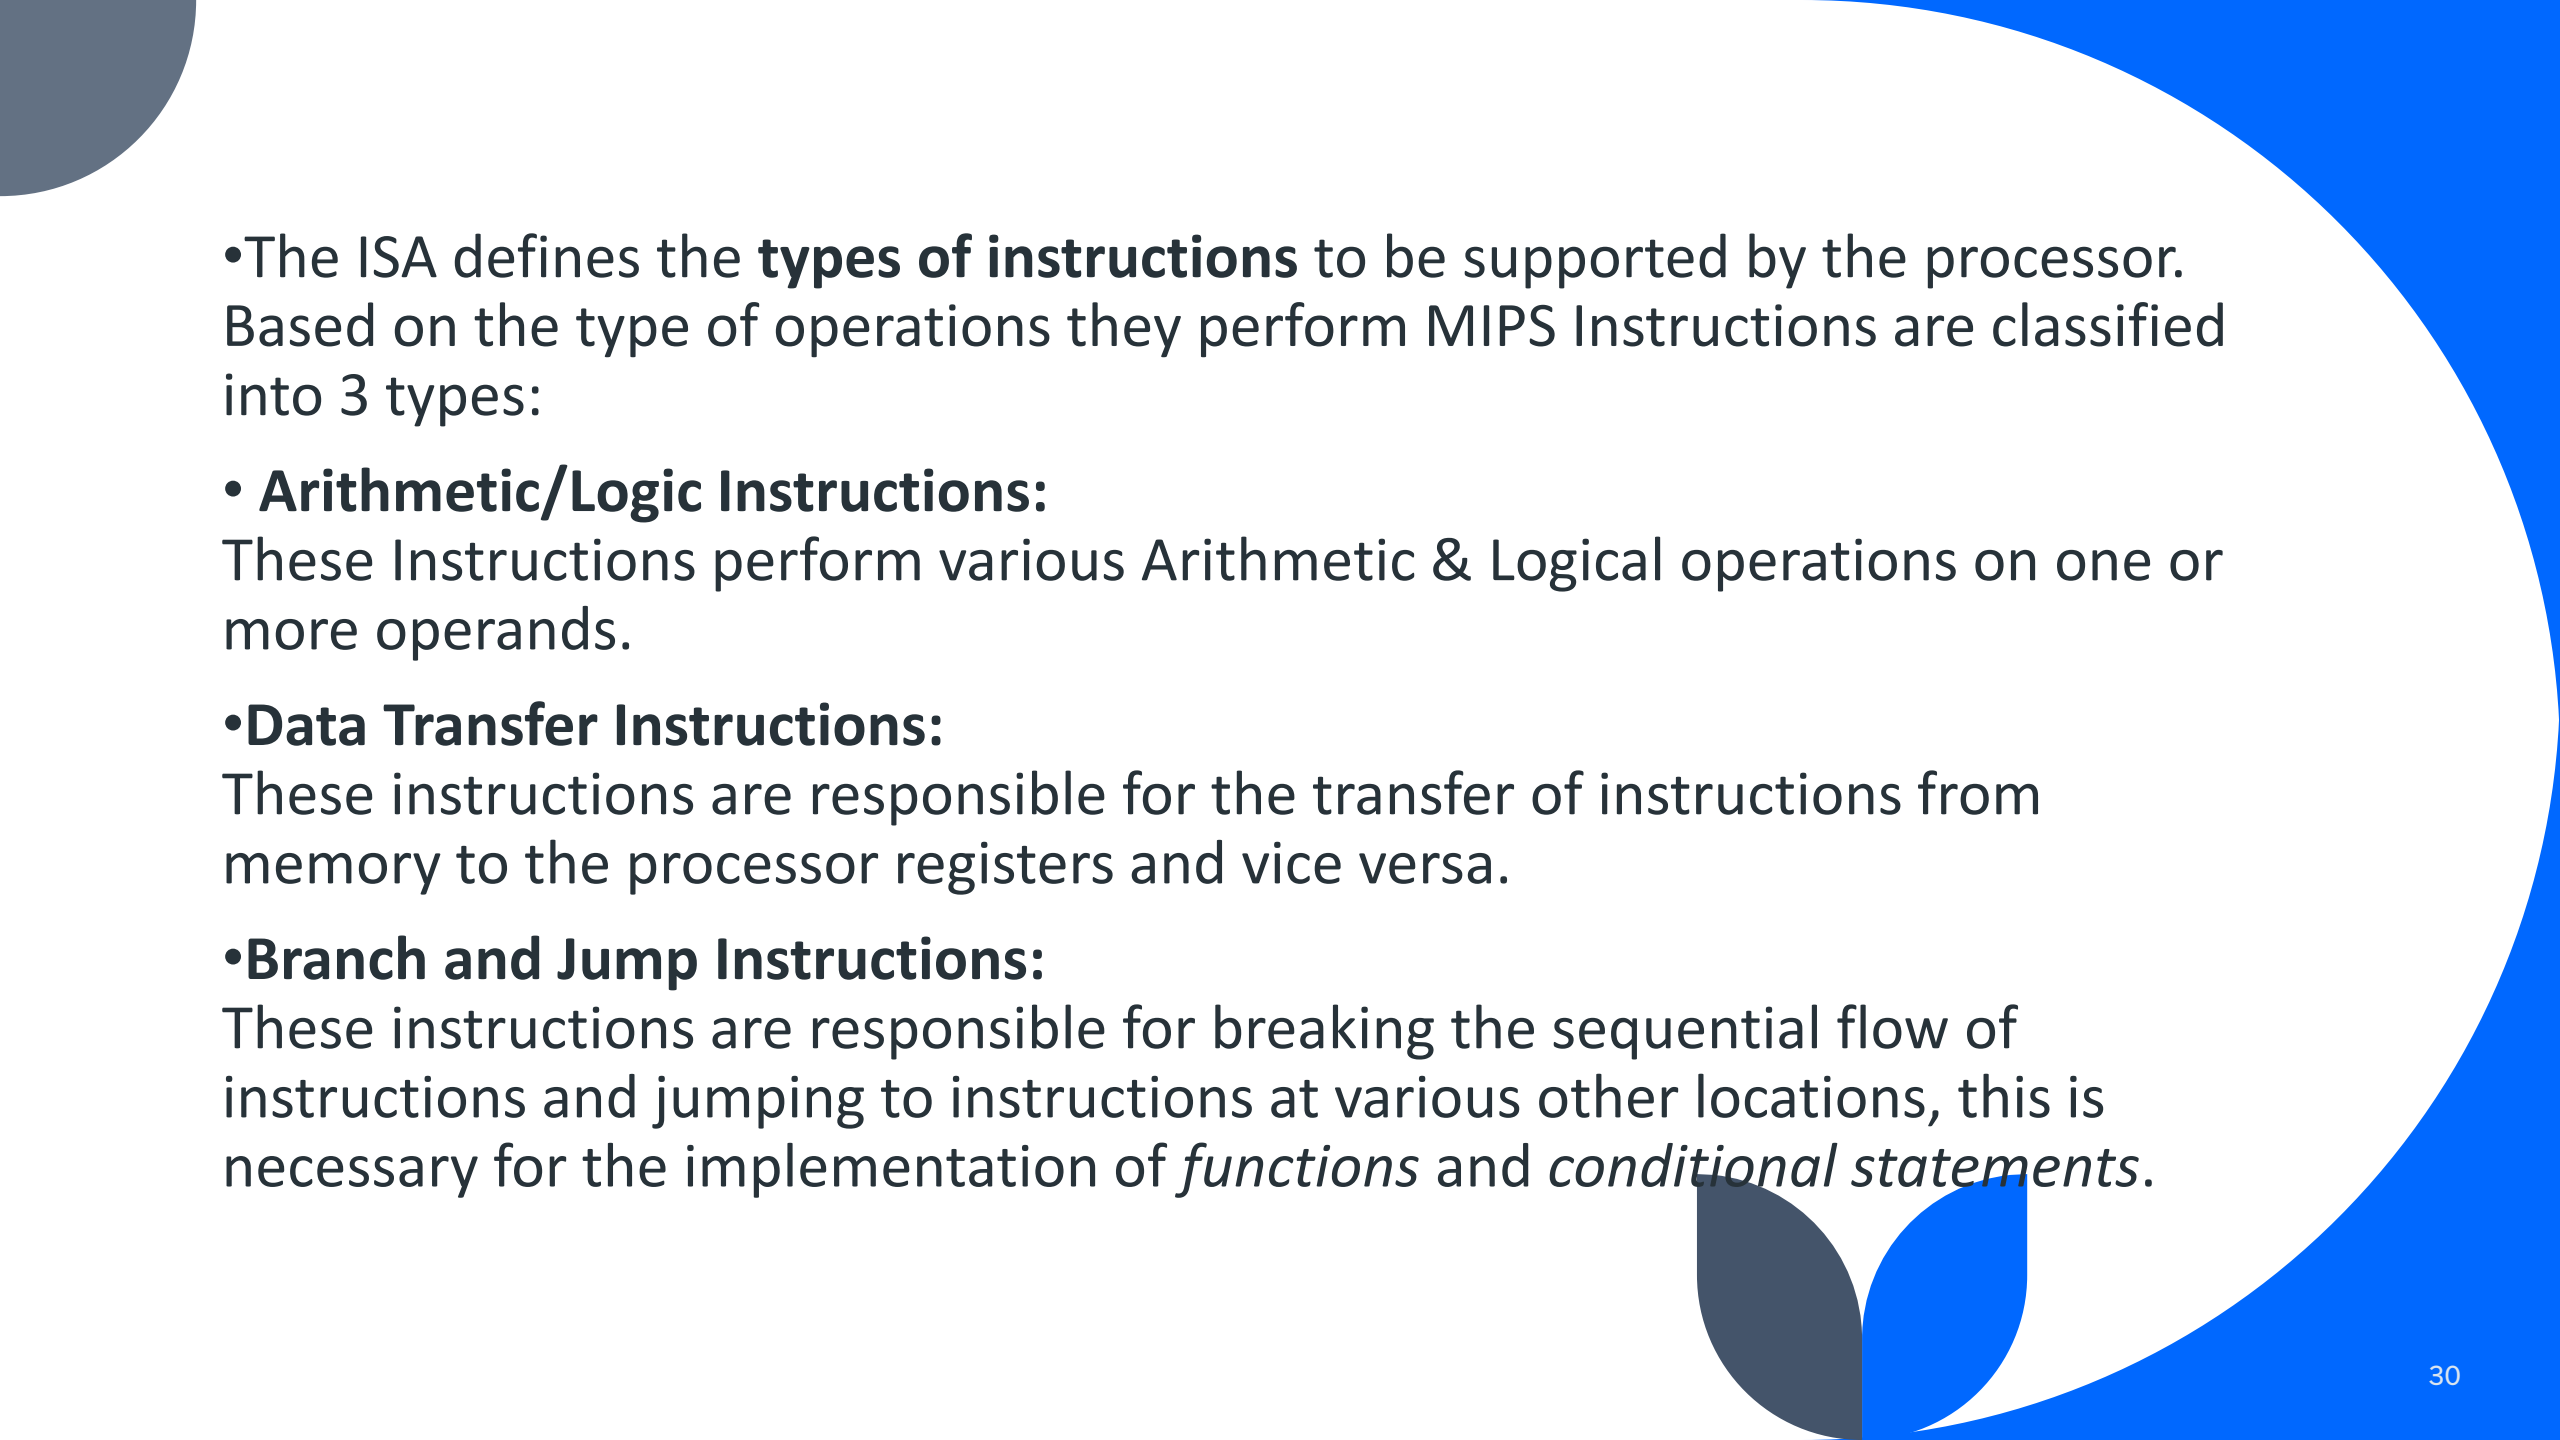
ISA(Instruction set Architecture)

In computer science, an instruction set architecture (ISA), also called computer architecture, is **an abstract model of a computer**. A device that executes instructions described by that ISA, such as a central processing unit (CPU), is called an implementation.



ISA describes the **design of a Computer** in terms of the **basic operations** it must support. The ISA is not concerned with the implementation-specific details of a computer. It is only concerned with the set or collection of basic operations the computer must support.





- The ISA defines the **types of instructions** to be supported by the processor. Based on the type of operations they perform MIPS Instructions are classified into 3 types:

- **Arithmetic/Logic Instructions:**

These Instructions perform various Arithmetic & Logical operations on one or more operands.

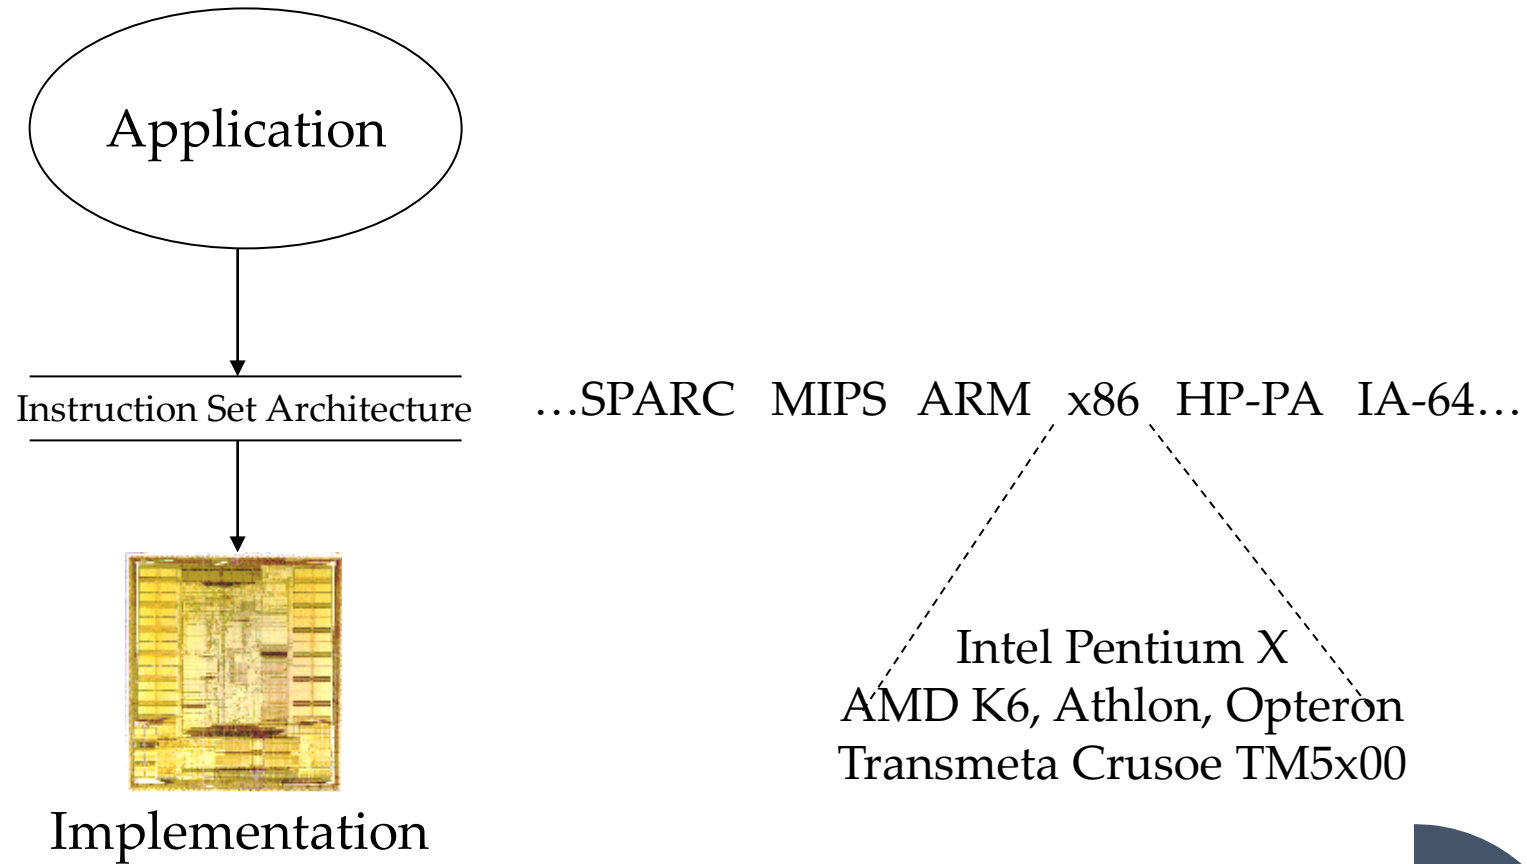
- **Data Transfer Instructions:**

These instructions are responsible for the transfer of instructions from memory to the processor registers and vice versa.

- **Branch and Jump Instructions:**

These instructions are responsible for breaking the sequential flow of instructions and jumping to instructions at various other locations, this is necessary for the implementation of *functions* and *conditional statements*.

Instruction Set Architecture



Instruction Set Architecture

Strong influence on cost/performance

New ISAs are rare, but versions are not
16-bit, 32-bit and 64-bit X86 versions

Longevity is a strong function of marketing
prowess

Traditional Issues

Strongly constrained by the number of bits available to instruction encoding

Opcodes/operands

Registers/memory

Addressing modes

Orthogonality

0, 1, 2, 3 address machines


Instruction formats

Decoding uniformity

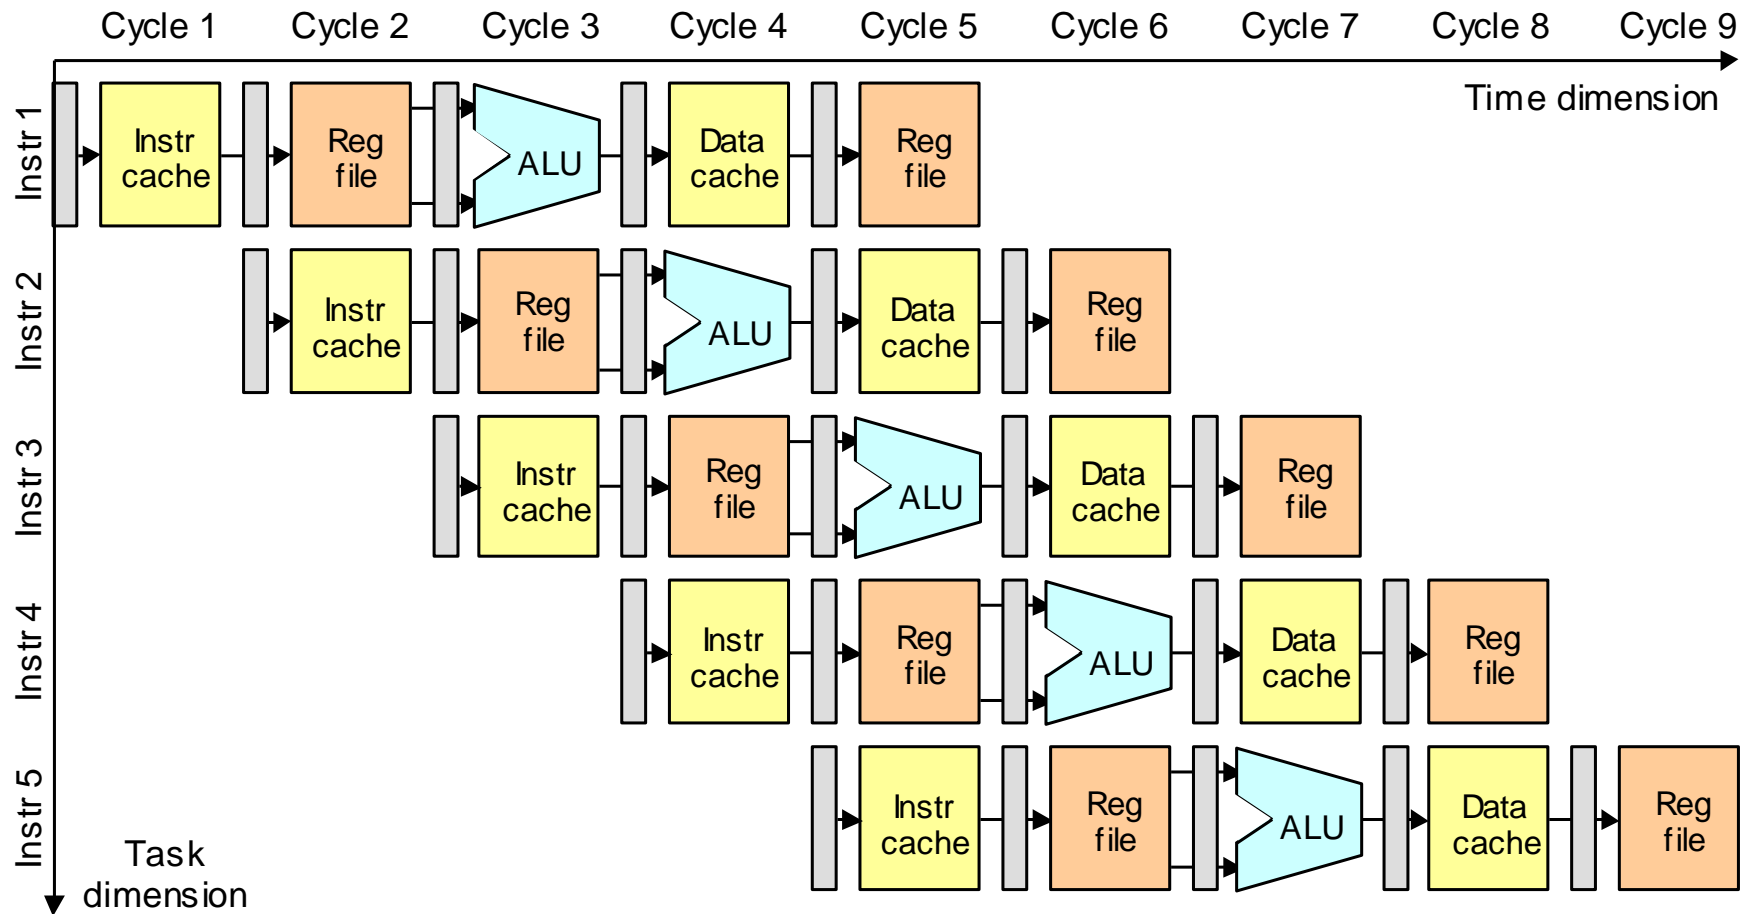
Pipelining Concepts

Strategies for improving performance

- 1 – Use multiple independent data paths accepting several instructions that are read out at once: *multiple-instruction-issue* or *superscalar*
- 2 – Overlap execution of several instructions, starting the next instruction before the previous one has run to completion: *(super)pipelined*

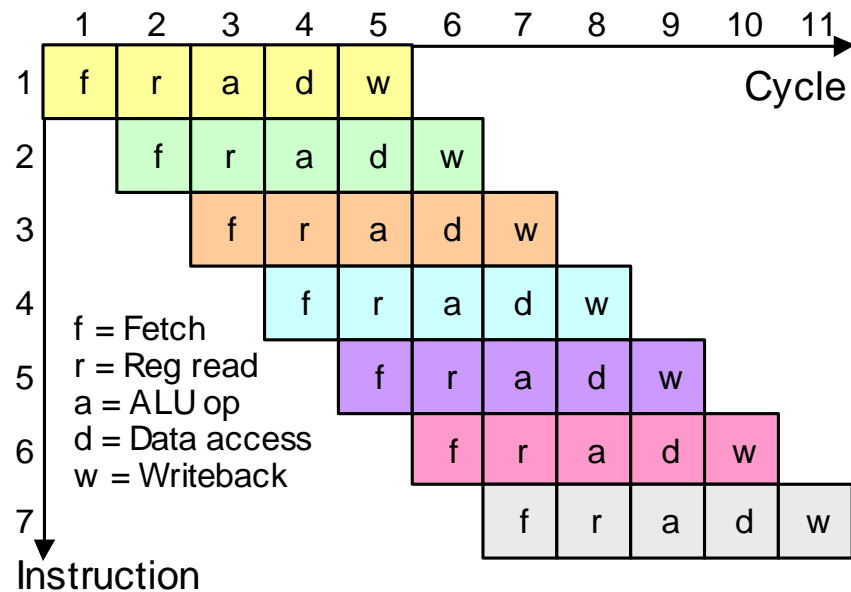


Pipelined Instruction Execution




Alternate Representations of a Pipeline

Except for start-up and drainage overheads, a pipeline can execute one instruction per clock tick; IPS is dictated by the clock frequency



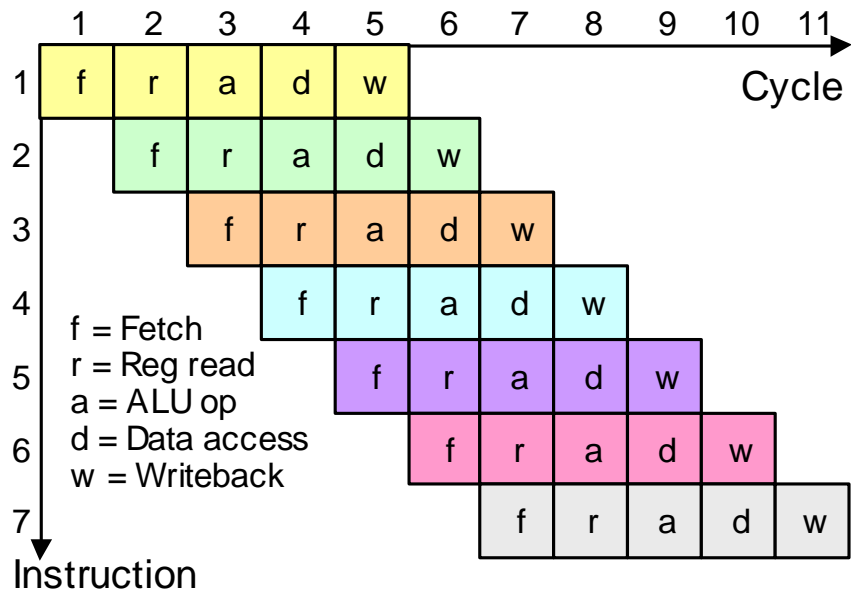
(a) Task-time diagram




(b) Space-time diagram

Alternate Representations of a Pipeline

Except for start-up and drainage overheads, a pipeline can execute one instruction per clock tick; IPS is dictated by the clock frequency



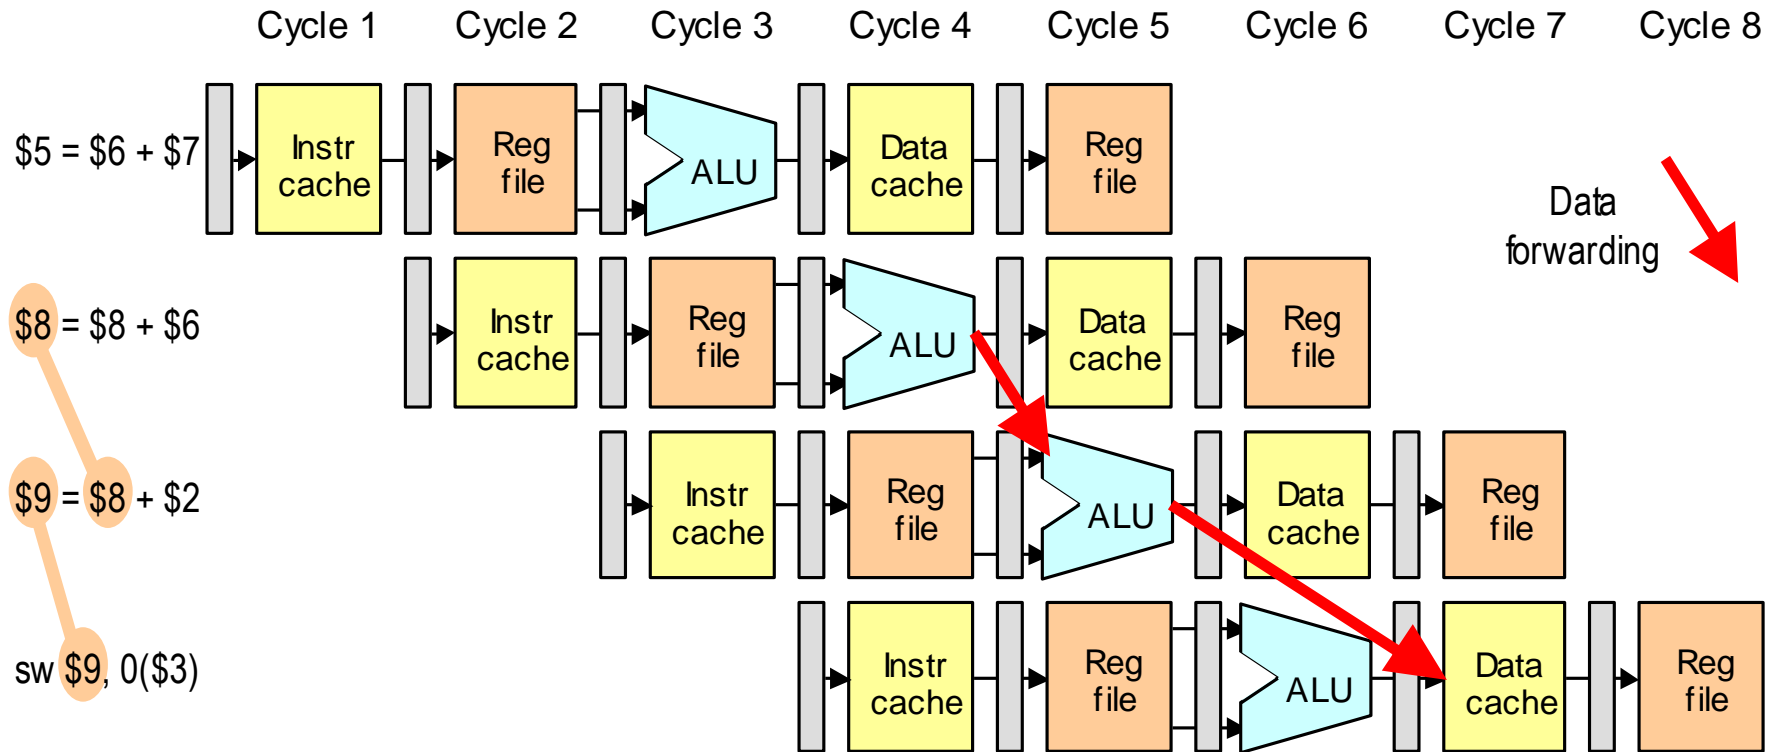
(a) Task-time diagram




(b) Space-time diagram





15.2 Pipeline Stalls or Bubbles

First type of data dependency






Writes into \$8





Reads from \$8

Without data forwarding,
three bubbles are needed
to resolve a read-after-write
data dependency



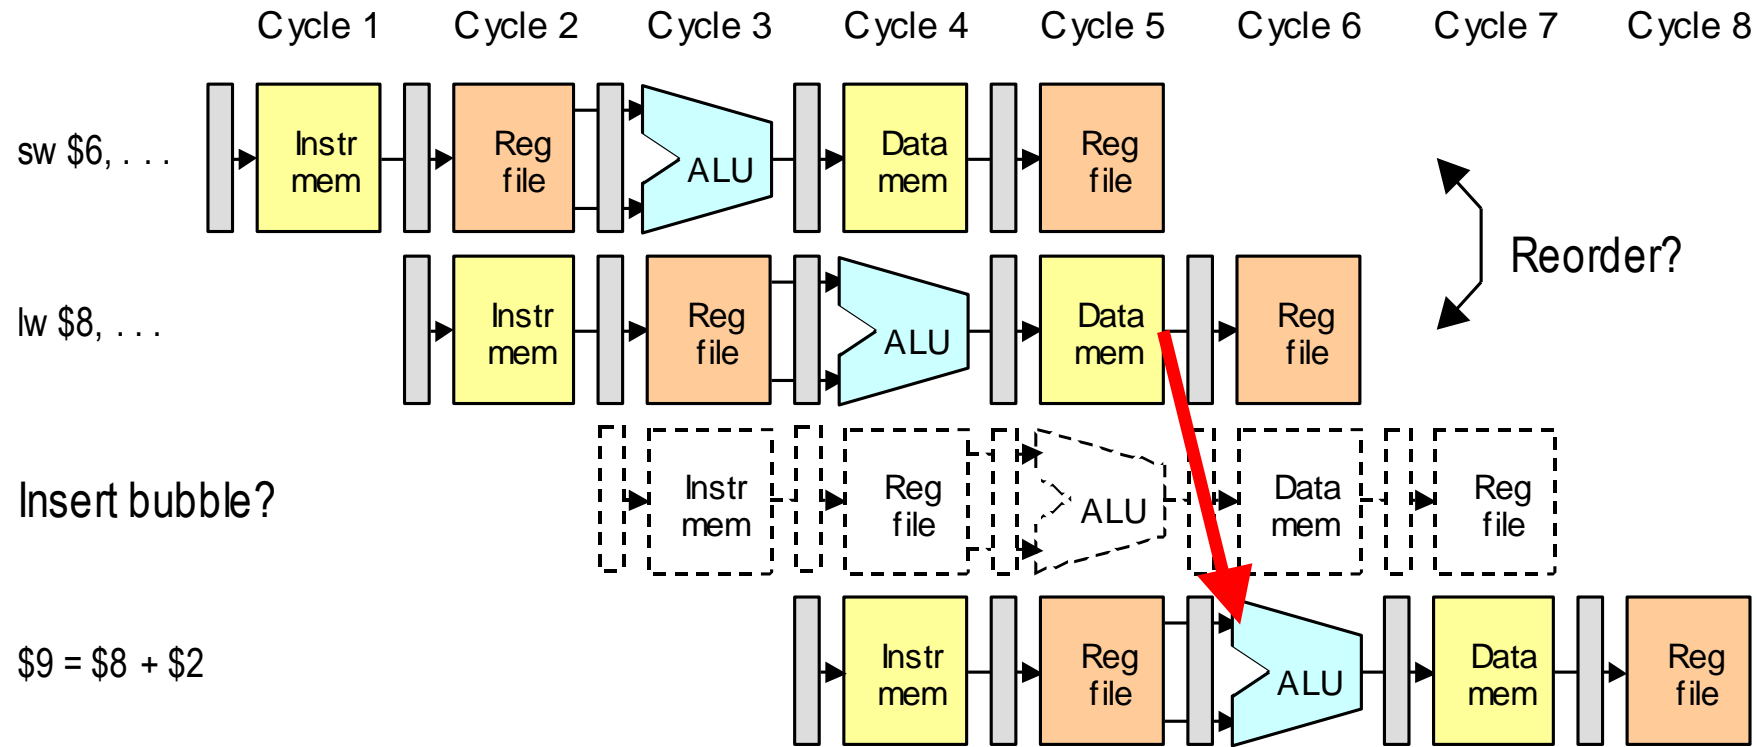
Writes into \$8



Reads from \$8

Two bubbles, if we assume
that a register can be updated
and read from in one cycle

Second Type of Data Dependency



Without data forwarding, three (two) bubbles are needed to resolve a read-after-load data dependency

Read-after-load data dependency and its possible resolution through bubble insertion and data forwarding.

Introduction

A.1 What is Pipelining?

A.2 The Major Hurdle of Pipelining-Structural Hazards

Data Hazards

Control Hazards

A.3 How is Pipelining Implemented

A.4 What Makes Pipelining Hard to Implement?

A.5 Extending the MIPS Pipeline to Handle Multi-cycle Operations

What Is Pipelining





Laundry Example

Ann, Brian, Cathy, Dave
each have one load of clothes
to wash, dry, and fold

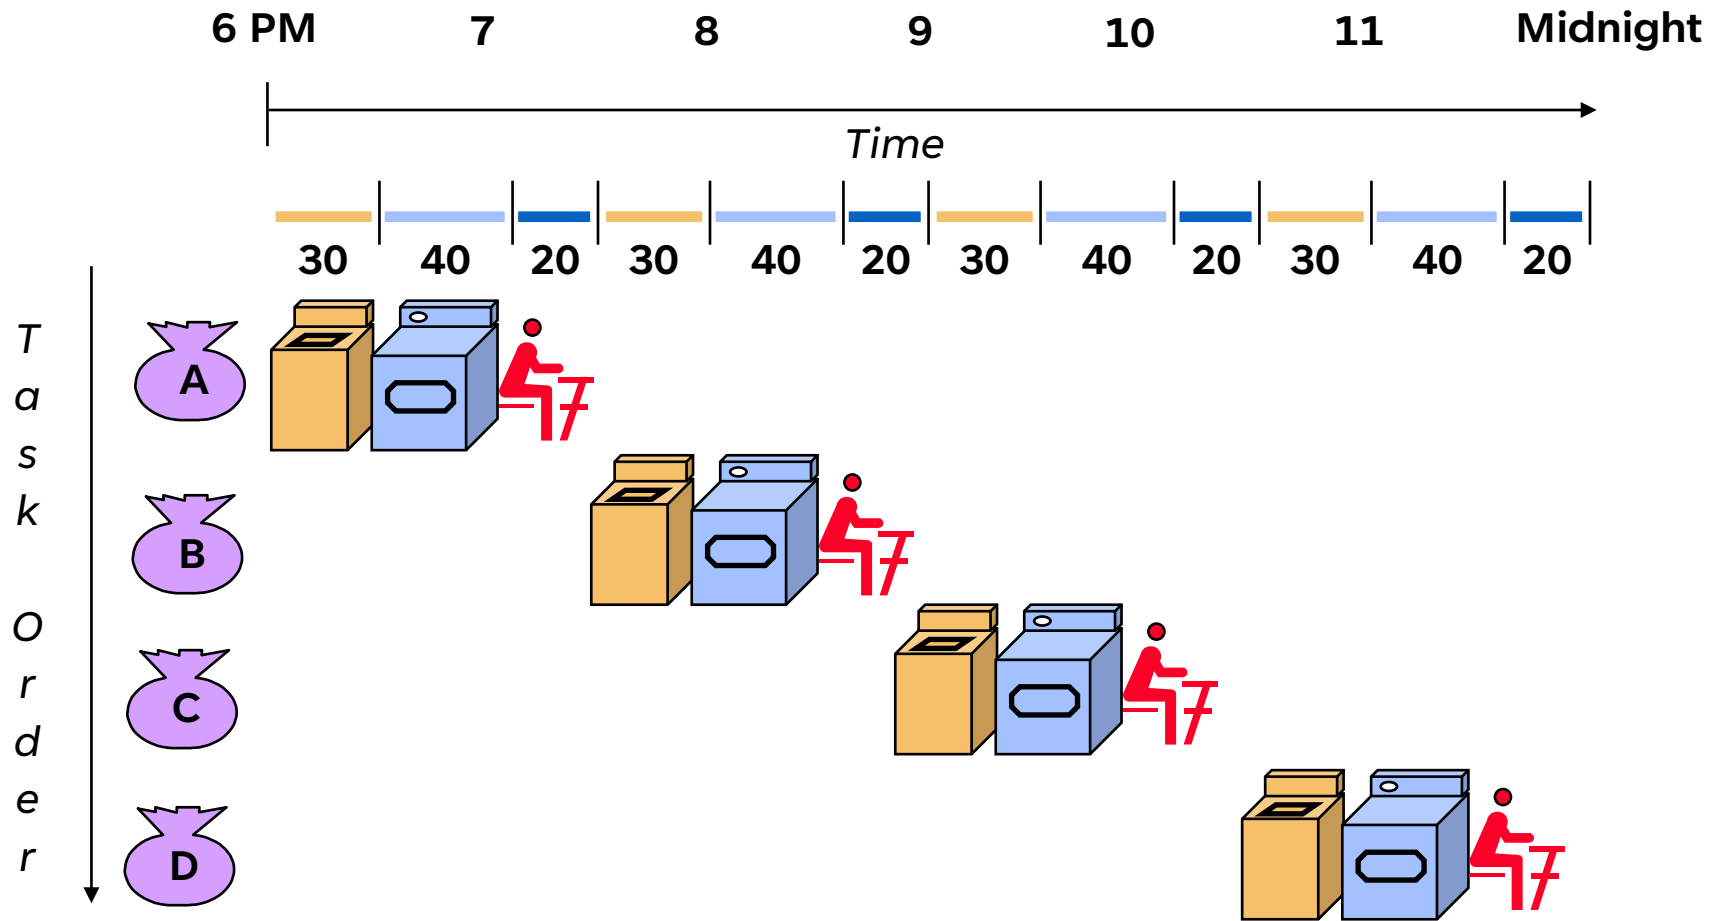
Washer takes 30 minutes

Dryer takes 40 minutes

“Folder” takes 20 minutes



What Is Pipelining




Sequential laundry takes 6 hours for 4 loads

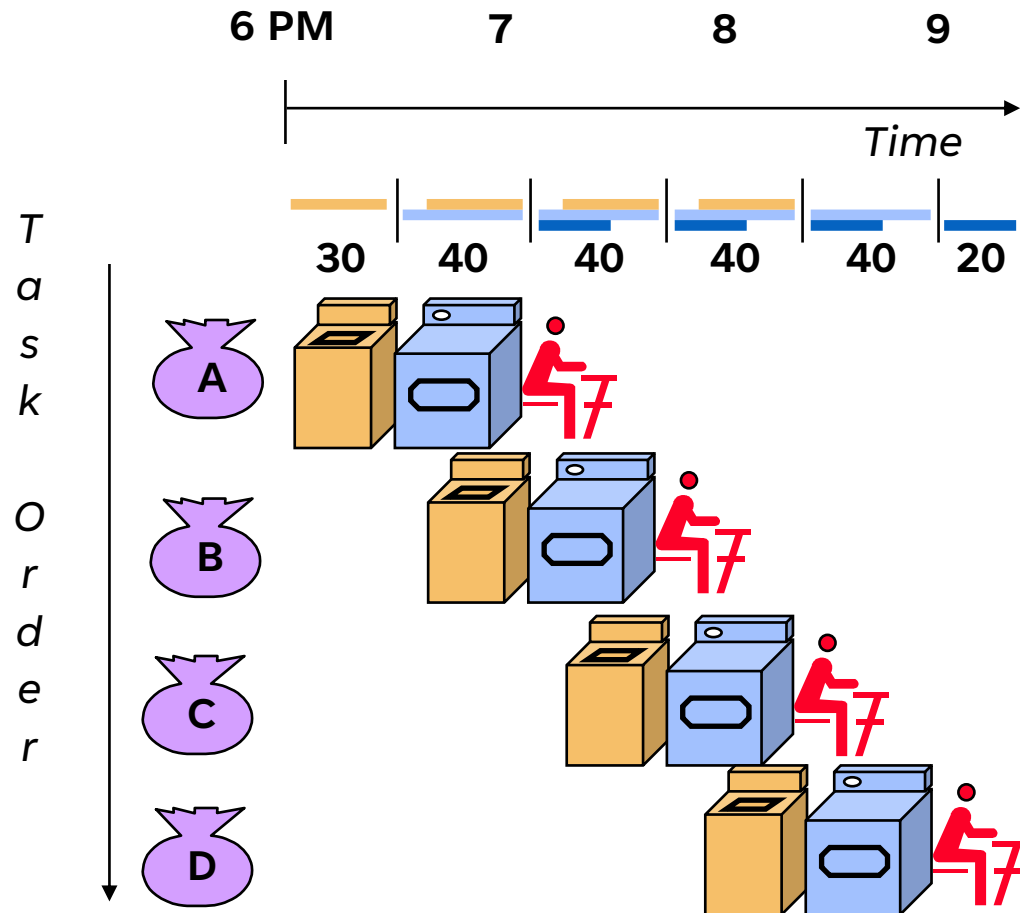
If they learned pipelining, how long would laundry take?

What Is Pipelining

Start work ASAP



What Is Pipelining



Pipelining Lessons

Pipelining doesn't help latency of single task, it helps throughput of entire workload

Pipeline rate limited by slowest pipeline stage

Multiple tasks operating simultaneously

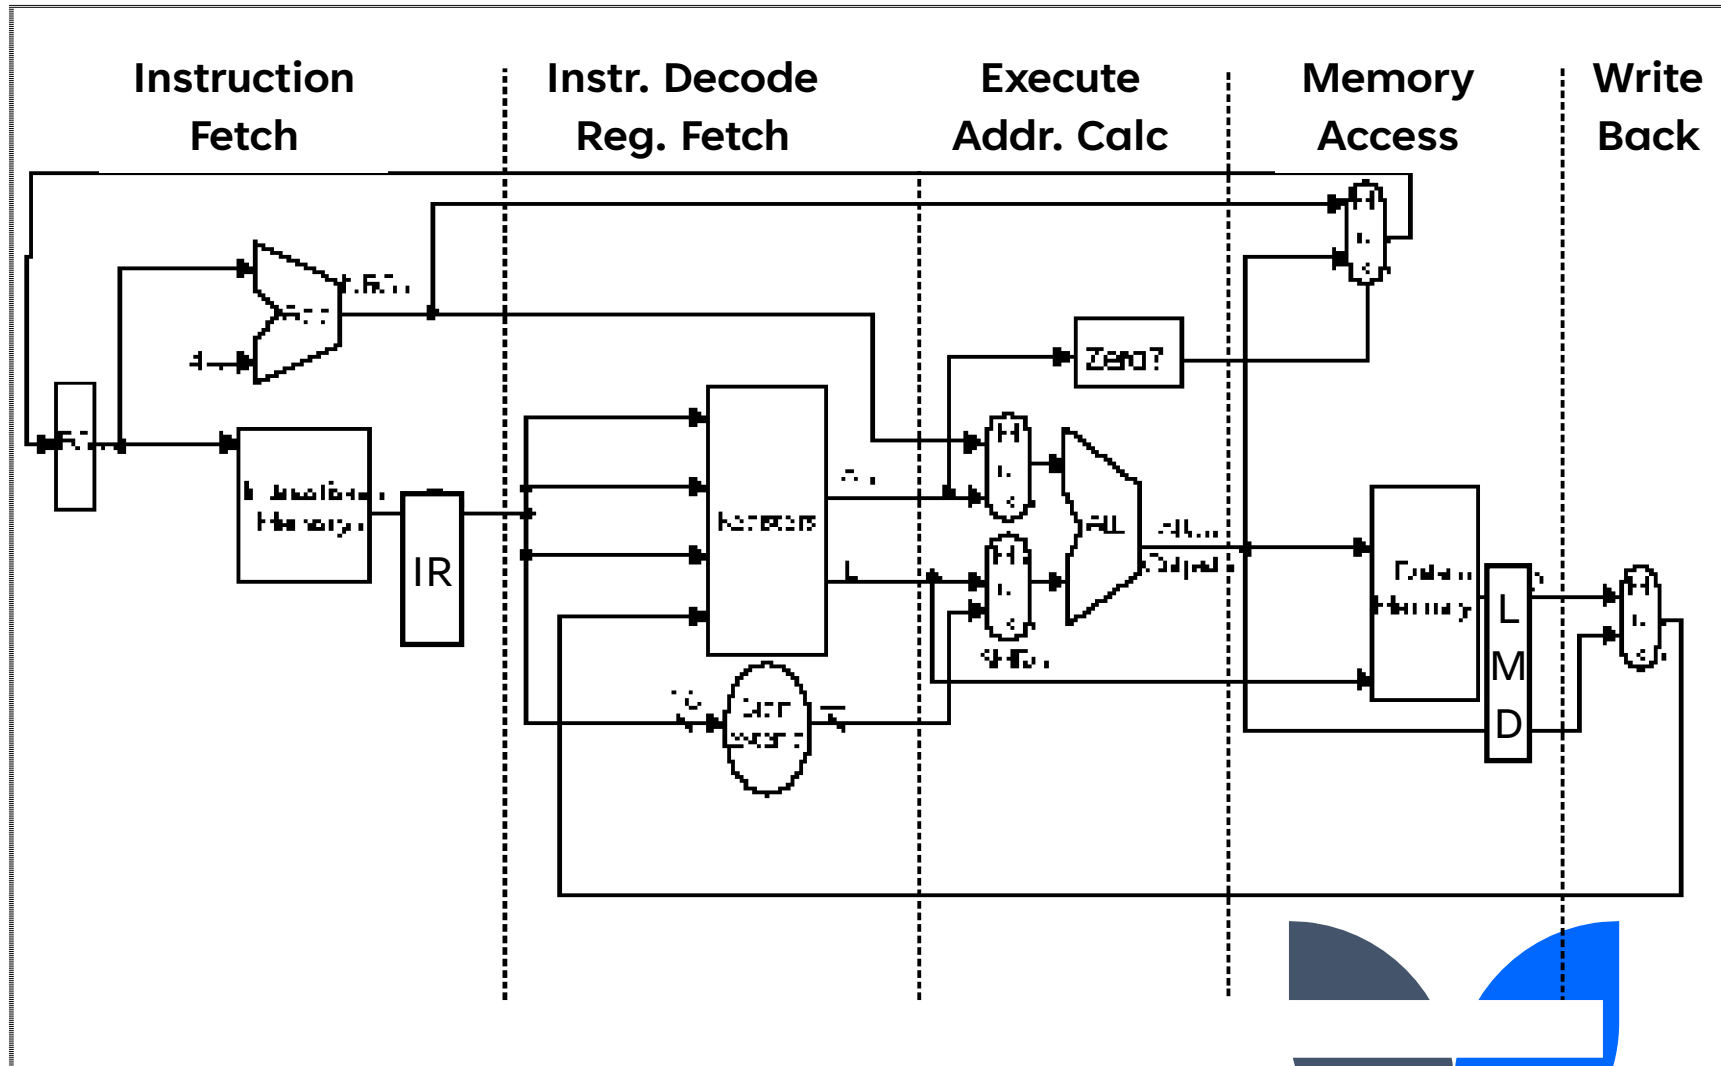
Potential speedup = Number pipe stages

Unbalanced lengths of pipe stages reduces speedup

Time to "fill" pipeline and time to "drain" it reduces speedup

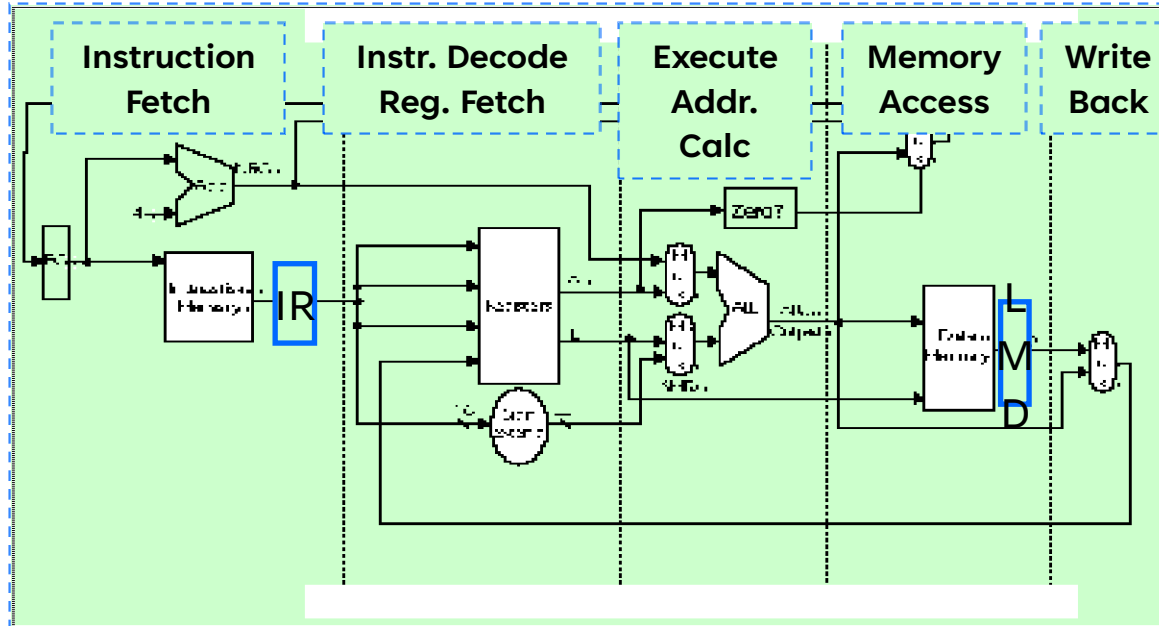
What Is Pipelining

MIPS Without Pipelining



What Is Pipelining

MIPS Functions



Passed To Next Stage

$IR \leftarrow \text{Mem}[PC]$
 $NPC \leftarrow PC + 4$

Instruction Fetch (IF):


Send out the PC and fetch the instruction from memory into the instruction register (IR); increment the PC by 4 to address the next sequential instruction.

IR holds the instruction that will be used in the next stage.

NPC holds the value of the next PC.

What Is Pipelining

MIPS Functions



Passed To Next Stage

```
A <- Regs[IR6..IR10];  
B <- Regs[IR10..IR15];  
Imm <- ((IR16) ##IR16-31
```

Instruction Decode/Register Fetch Cycle (ID):

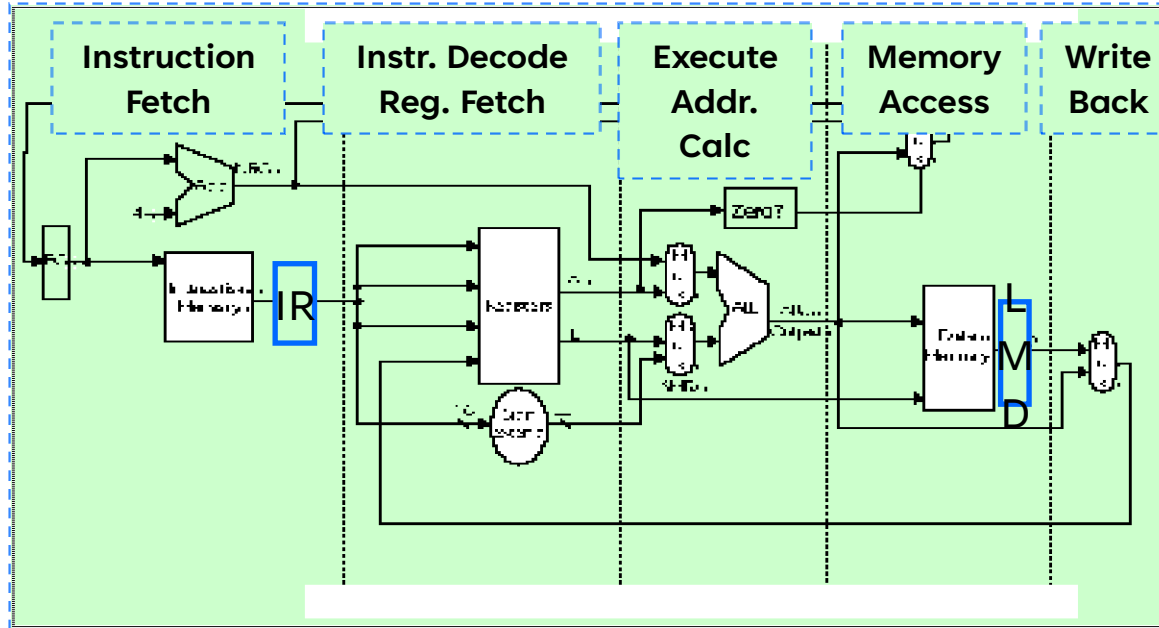
Decode the instruction and access the register file to read the registers.

The outputs of the general purpose registers are read into two temporary registers (A & B) for use in later clock cycles.

We extend the sign of the lower 16 bits of the Instruction Register.

What Is Pipelining

MIPS Functions



Passed To Next Stage

`A <- A func. B`
`cond = 0;`

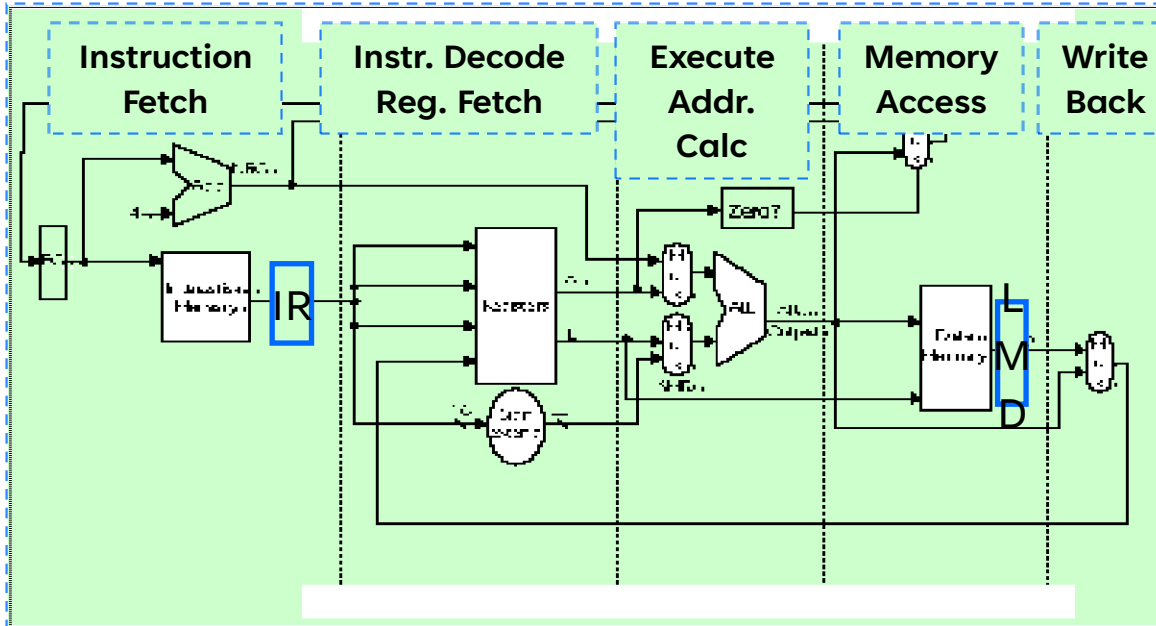
Execute Address Calculation (EX):

We perform an operation (for an ALU) or an address calculation (if it's a load or a Branch).

If an ALU, actually do the operation. If an address calculation, figure out how to obtain the address and stash away the location of that address for the next cycle.

What Is Pipelining

MIPS Functions



Passed To Next Stage

$A = \text{Mem}[\text{prev. B}]$
or
 $\text{Mem}[\text{prev. B}] = A$


MEMORY ACCESS (MEM):

If this is an ALU, do nothing.

If a load or store, then access memory.

What Is Pipelining

MIPS Functions

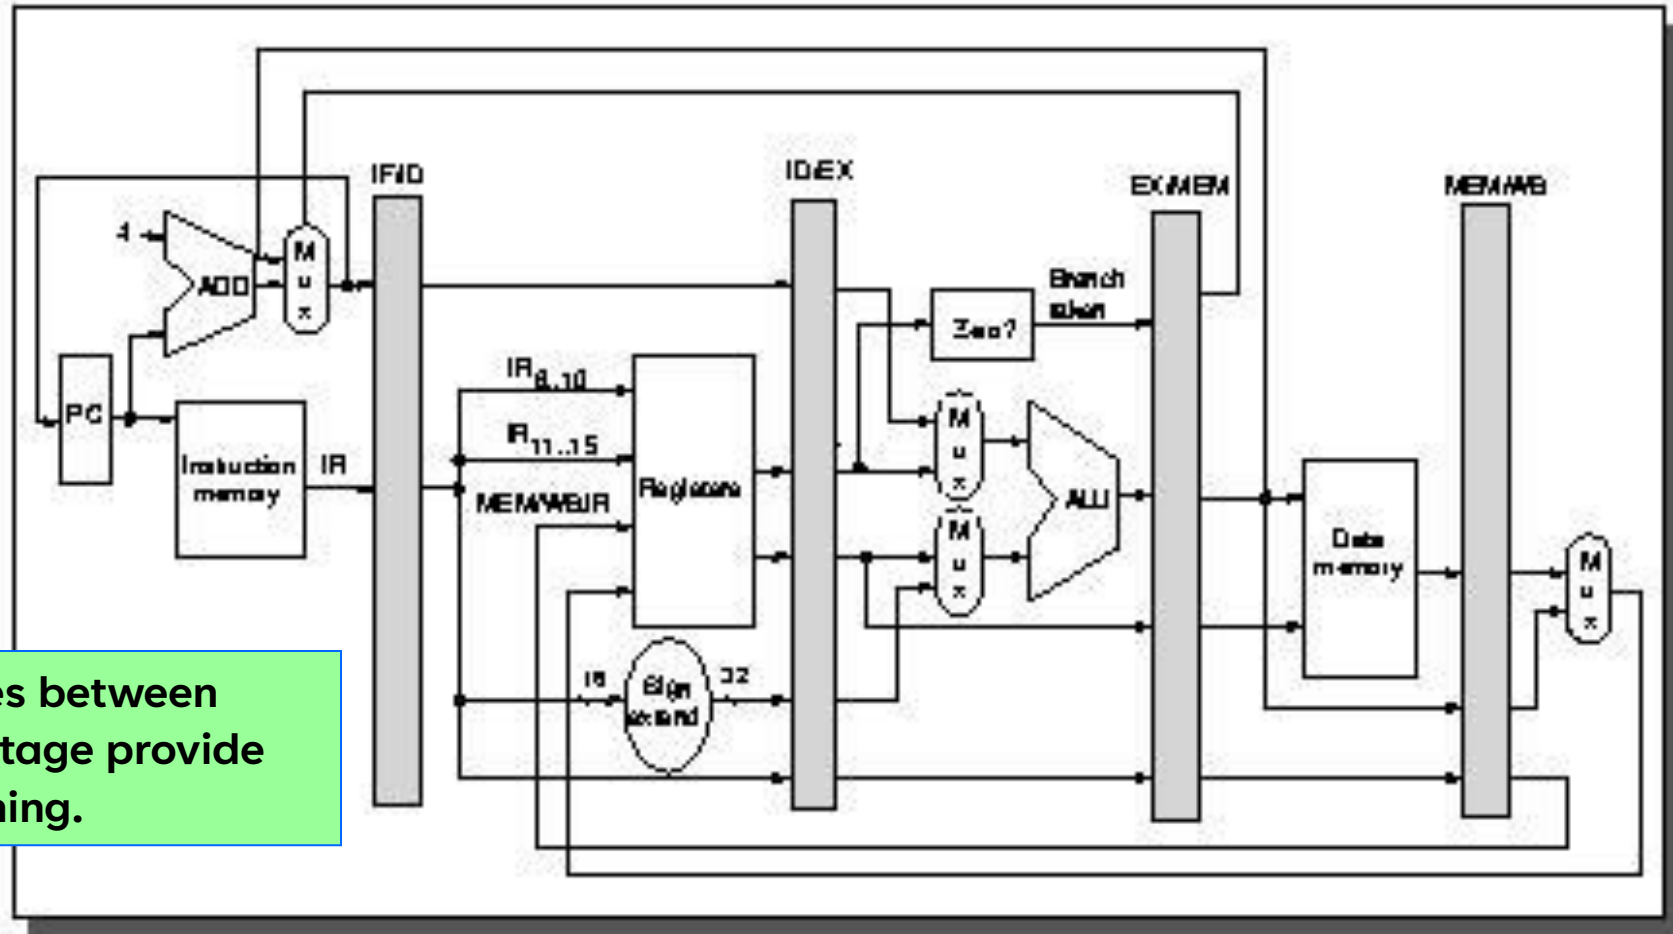


Passed To Next Stage
Regs \leftarrow A, B;

WRITE BACK (WB):

Update the registers from either the ALU or from the data loaded.

The Basic Pipeline For MIPS



Latches between each stage provide pipelining.

FIGURE 3.4 The datapath is pipelined by adding a set of registers, one between each pair of pipe stages.

The Basic Pipeline For MIPS

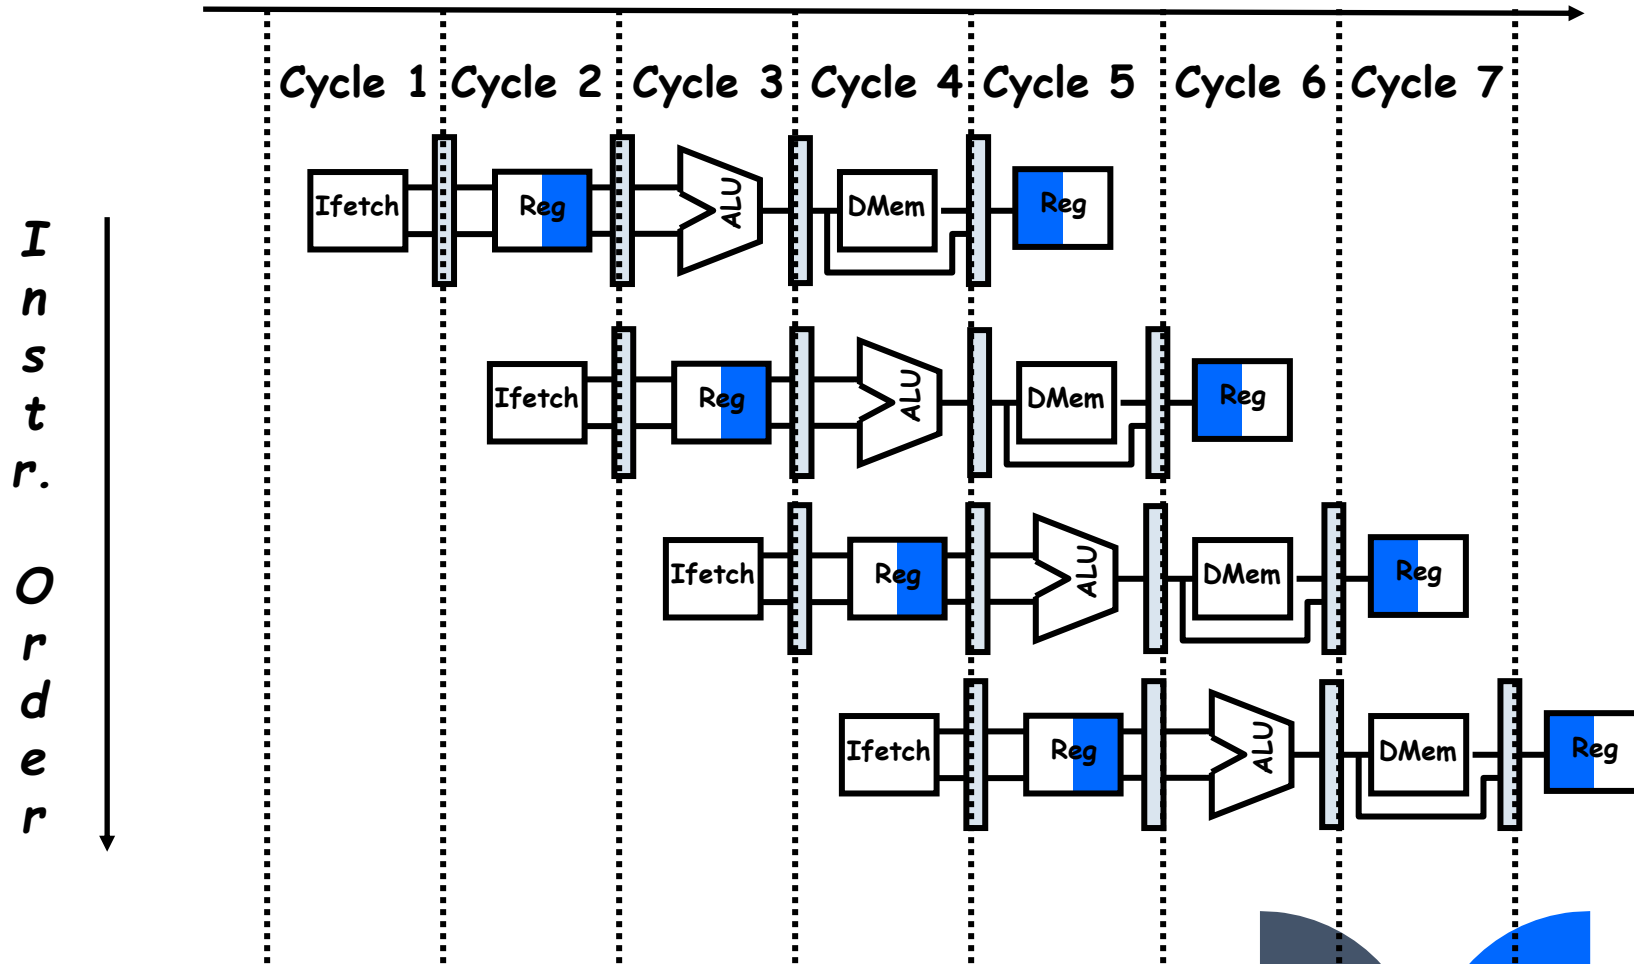


Figure 3.3

Pipeline Hurdles

A.1 What is Pipelining?

A.2 The Major Hurdle of Pipelining- Structural Hazards

- Structural Hazards
- Data Hazards
- Control Hazards

A.3 How is Pipelining Implemented

A.4 What Makes Pipelining Hard to Implement?

A.5 Extending the MIPS Pipeline to Handle Multi-cycle Operations

Limits to pipelining: Hazards prevent next instruction from executing during its designated clock cycle

Structural hazards: HW cannot support this combination of instructions (single person to fold and put clothes away)

Data hazards: Instruction depends on result of prior instruction still in the pipeline (missing sock)

Control hazards: Pipelining of branches & other instructions that change the PC

Common solution is to **stall** the pipeline until the hazard is resolved, inserting one or more “**bubbles**” in the pipeline

Pipeline Hurdles

Definition

conditions that lead to incorrect behavior if not fixed

Structural hazard

two different instructions use same h/w in same cycle

Data hazard

two different instructions use same storage

must appear as if the instructions execute in correct order

Control hazard

one instruction affects which instruction is next

Resolution

Pipeline interlock logic detects hazards and fixes them

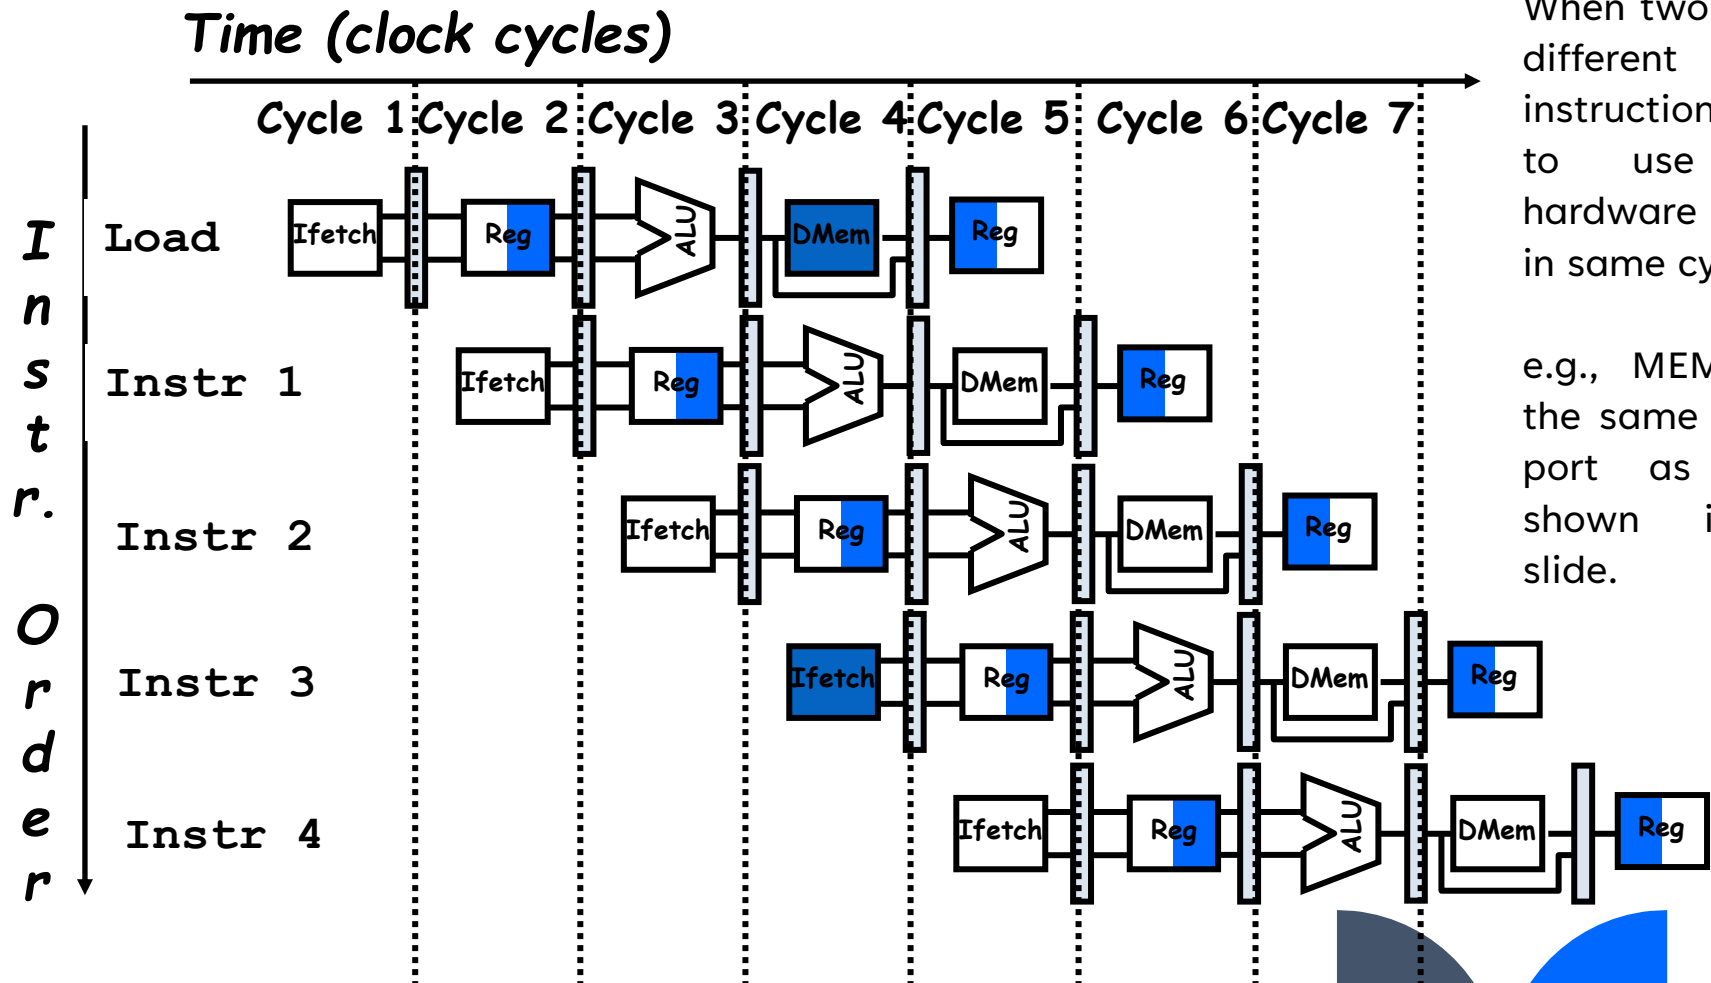
simple solution: stall

increases CPI, decreases performance

better solution: partial stall

some instruction stall, others proceed better to stall early than late

Structural Hazards

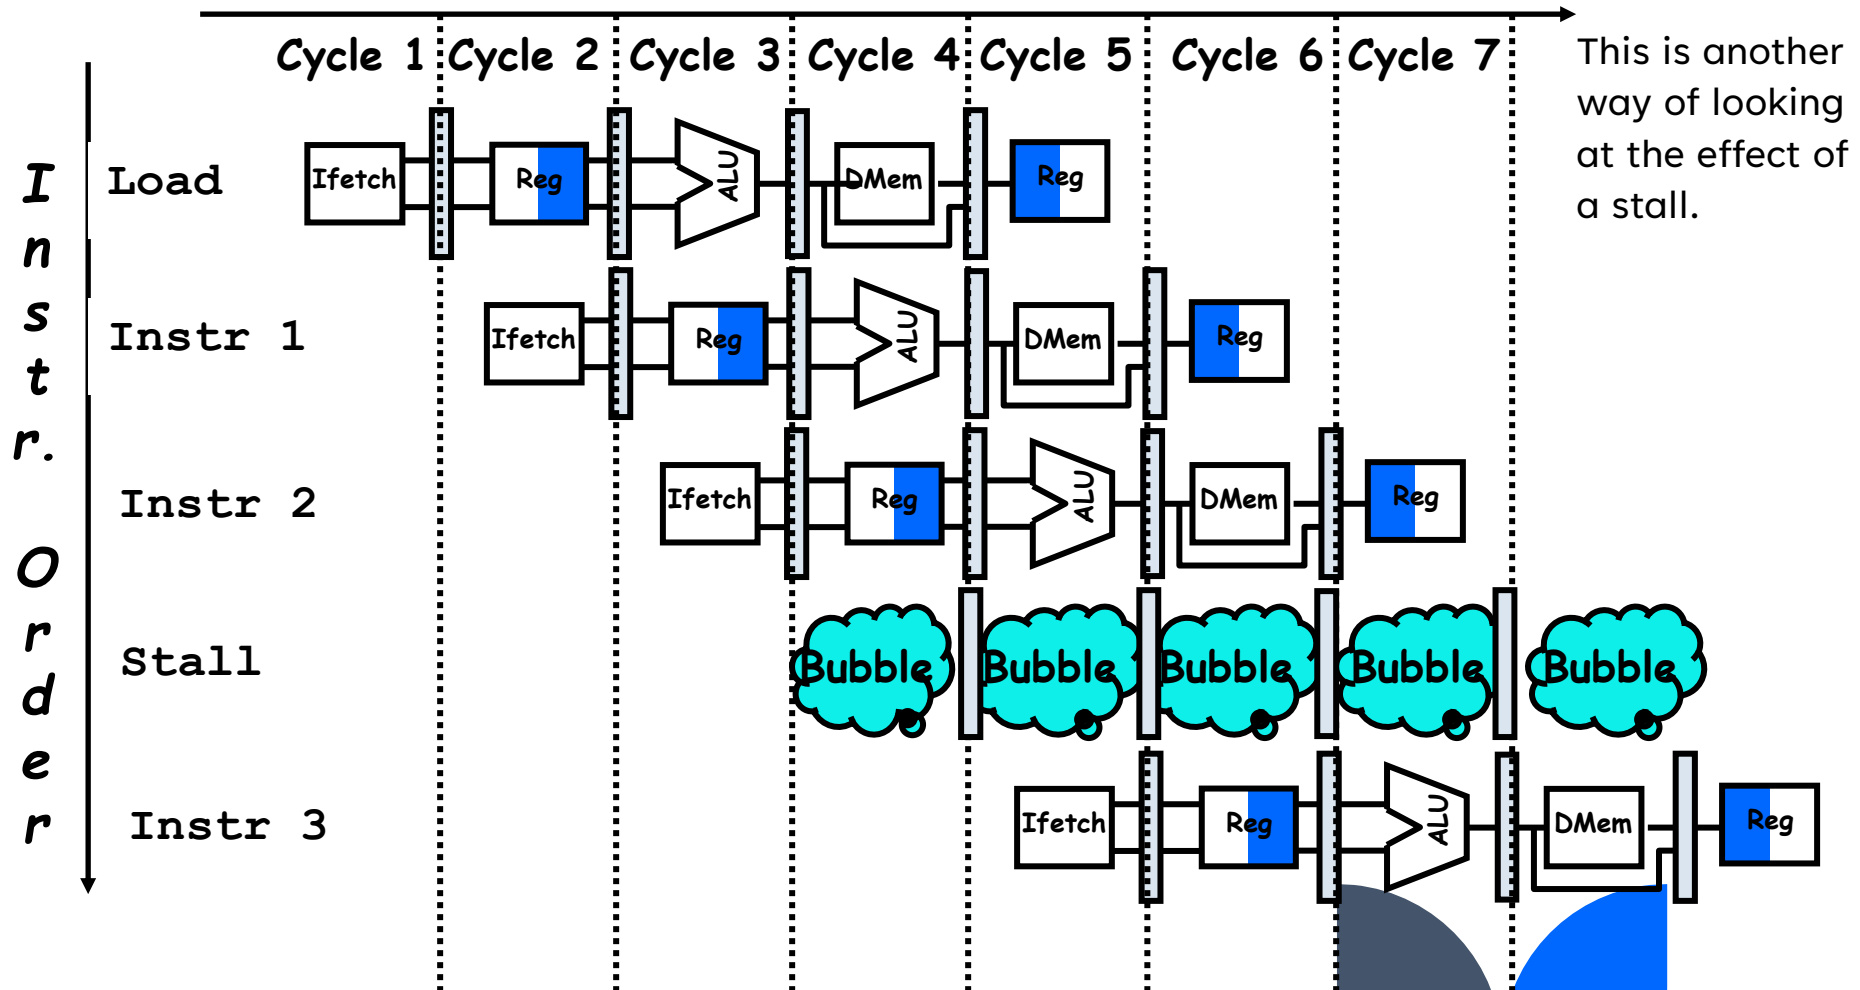


When two or more different instructions want to use same hardware resource in same cycle

e.g., MEM uses the same memory port as IF as shown in this slide.

Structural Hazards

Time (clock cycles)



Structural Hazards

Instruction	Clock cycle number									
	1	2	3	4	5	6	7	8	9	10
Load instruction	IF	ID	EX	MEM	WB					
Instruction $j+1$		IF	ID	EX	MEM	WB				
Instruction $j+2$			IF	ID	EX	MEM	WB			
Instruction $j+3$				stall	IF	ID	EX	MEM	WB	
Instruction $j+4$						IF	ID	EX	MEM	WB
Instruction $j+5$							IF	ID	EX	MEM
Instruction $j+6$								IF	ID	EX

This is another way to represent the stall we saw on the last few pages.

Structural Hazards

Dealing with Structural Hazards

Stall

low cost, simple

Increases CPI

use for rare case since stalling has performance effect

Pipeline hardware resource

useful for multi-cycle resources

good performance

sometimes complex e.g., RAM

Replicate resource

good performance

increases cost (+ maybe interconnect delay)

useful for cheap or divisible resources

Structural Hazards

Structural hazards are reduced with these rules:

Each instruction uses a resource at most once

Always use the resource in the same pipeline stage

Use the resource for one cycle only

Many RISC ISA's designed with this in mind

Sometimes very complex to do this. For example, memory of necessity is used in the IF and MEM stages.

Some common Structural Hazards:

- **Memory** - we've already mentioned this one.
- **Floating point** - Since many floating point instructions require many cycles, it's easy for them to interfere with each other.
- **Starting up more of one type of instruction than there are resources.** For instance, the PA-8600 can support two ALU + two load/store instructions per cycle - that's how much hardware it has available.

Data Hazards

A.1 What is Pipelining?

A.2 The Major Hurdle of Pipelining- Structural Hazards

- Structural Hazards
- Data Hazards
- Control Hazards

A.3 How is Pipelining Implemented

A.4 What Makes Pipelining Hard to Implement?

A.5 Extending the MIPS Pipeline to Handle Multi-cycle Operations

These occur when at any time, there are instructions active that need to access the same data (memory or register) locations.

Where there's real trouble is when we have:

instruction A

instruction B

and B manipulates (reads or writes) data before A does. This violates the order of the instructions, since the architecture implies that A completes entirely before B is executed.

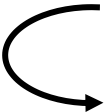
Data Hazards

Execution Order is:

Instr_i
Instr_j

Read After Write (RAW)

Instr_j tries to read operand before Instr_i writes it

 I: add **r1**, r2, r3
J: sub r4, **r1**, r3

Caused by a “**Dependence**” (in compiler nomenclature). This hazard results from an actual need for communication.

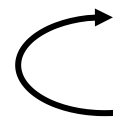
Data Hazards

Execution Order is:

Instr_i
Instr_j

Write After Read (WAR)

Instr_j tries to write operand before Instr_i reads it
Gets wrong operand



```
I: sub r4, r1, r3
J: add r1, r2, r3
K: mul r6, r1, r7
```

Called an “[anti-dependence](#)” by compiler writers.
This results from reuse of the name “[r1](#)”.

Can’t happen in MIPS 5 stage pipeline because:
All instructions take 5 stages, and
Reads are always in stage 2, and
Writes are always in stage 5

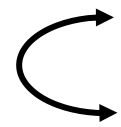
Data Hazards

Execution Order is:

Instr_i
Instr_j

Write After Write (WAW)

Instr_j tries to write operand before Instr_i writes it
Leaves wrong result (Instr_i not Instr_j)



```
I: sub r1, r4, r3
J: add r1, r2, r3
K: mul r6, r1, r7
```

Called an “**output dependence**” by compiler writers
This also results from the reuse of name “**r1**”.

Can’t happen in MIPS 5 stage pipeline because:
All instructions take 5 stages, and
Writes are always in stage 5

Will see WAR and WAW in later more complicated pipes

Data Hazards

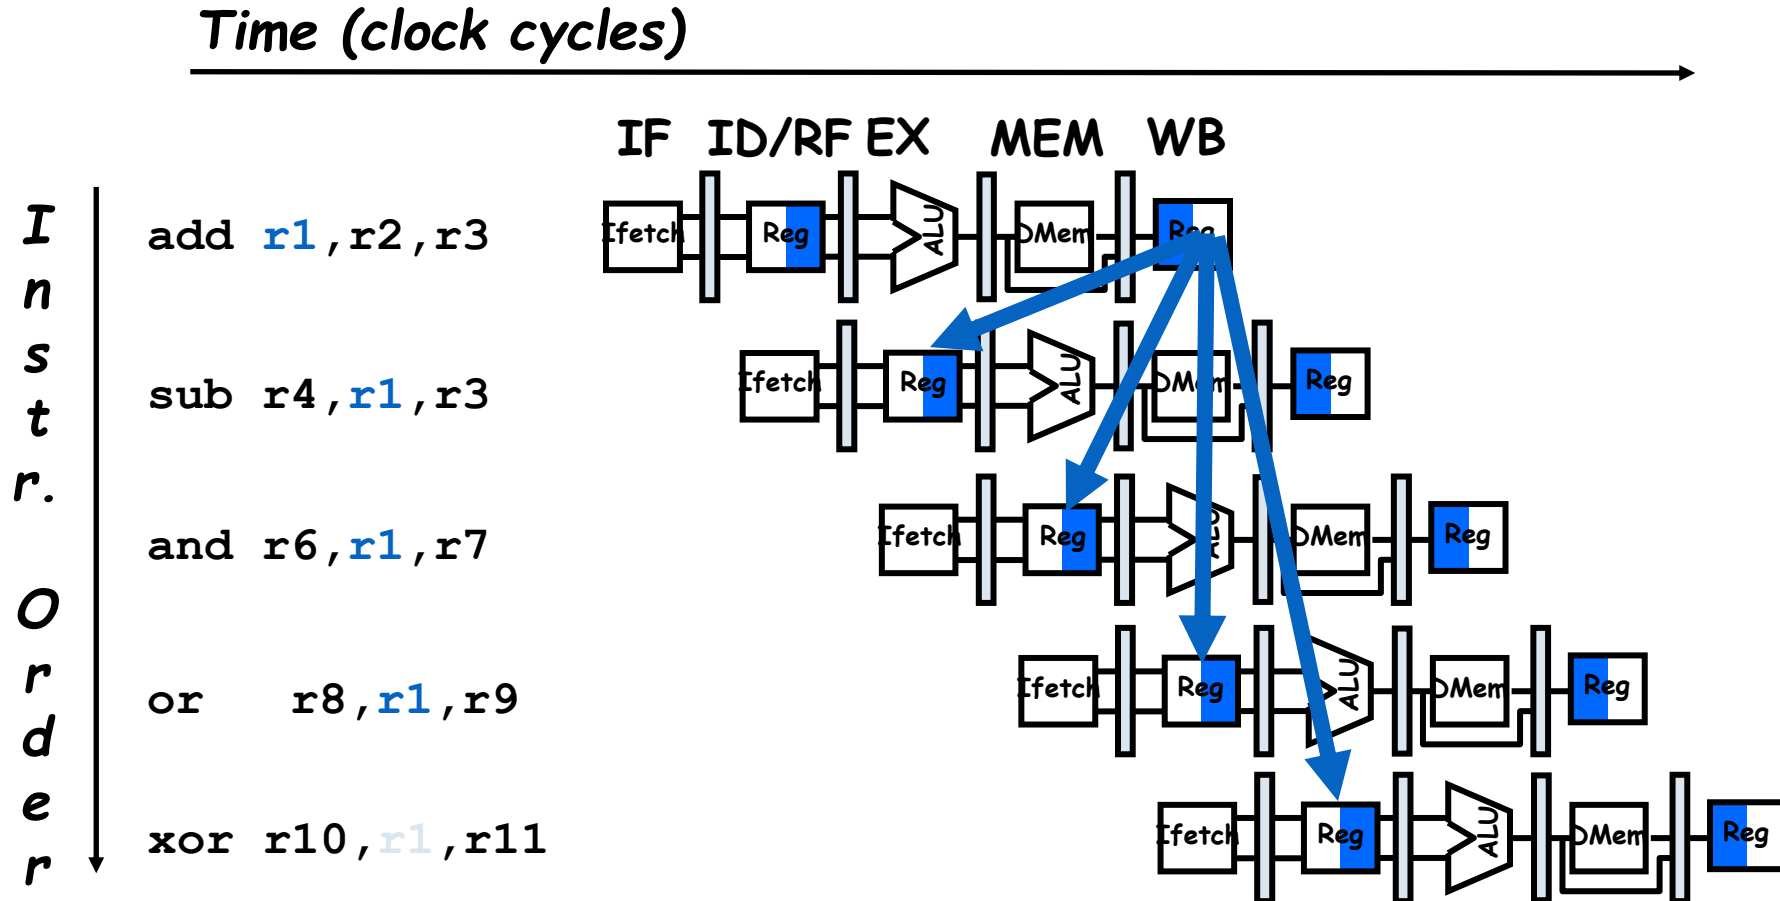
Simple Solution to RAW

- Hardware detects RAW and stalls
- Assumes register written then read each cycle
 - + low cost to implement, simple
 - reduces IPC
- Try to minimize stalls

Minimizing RAW stalls

- Bypass/forward/shortcircuit (We will use the word “forward”)
- Use data before it is in the register
 - + reduces/avoids stalls
 - complex
- Crucial for common RAW hazards

Data Hazards

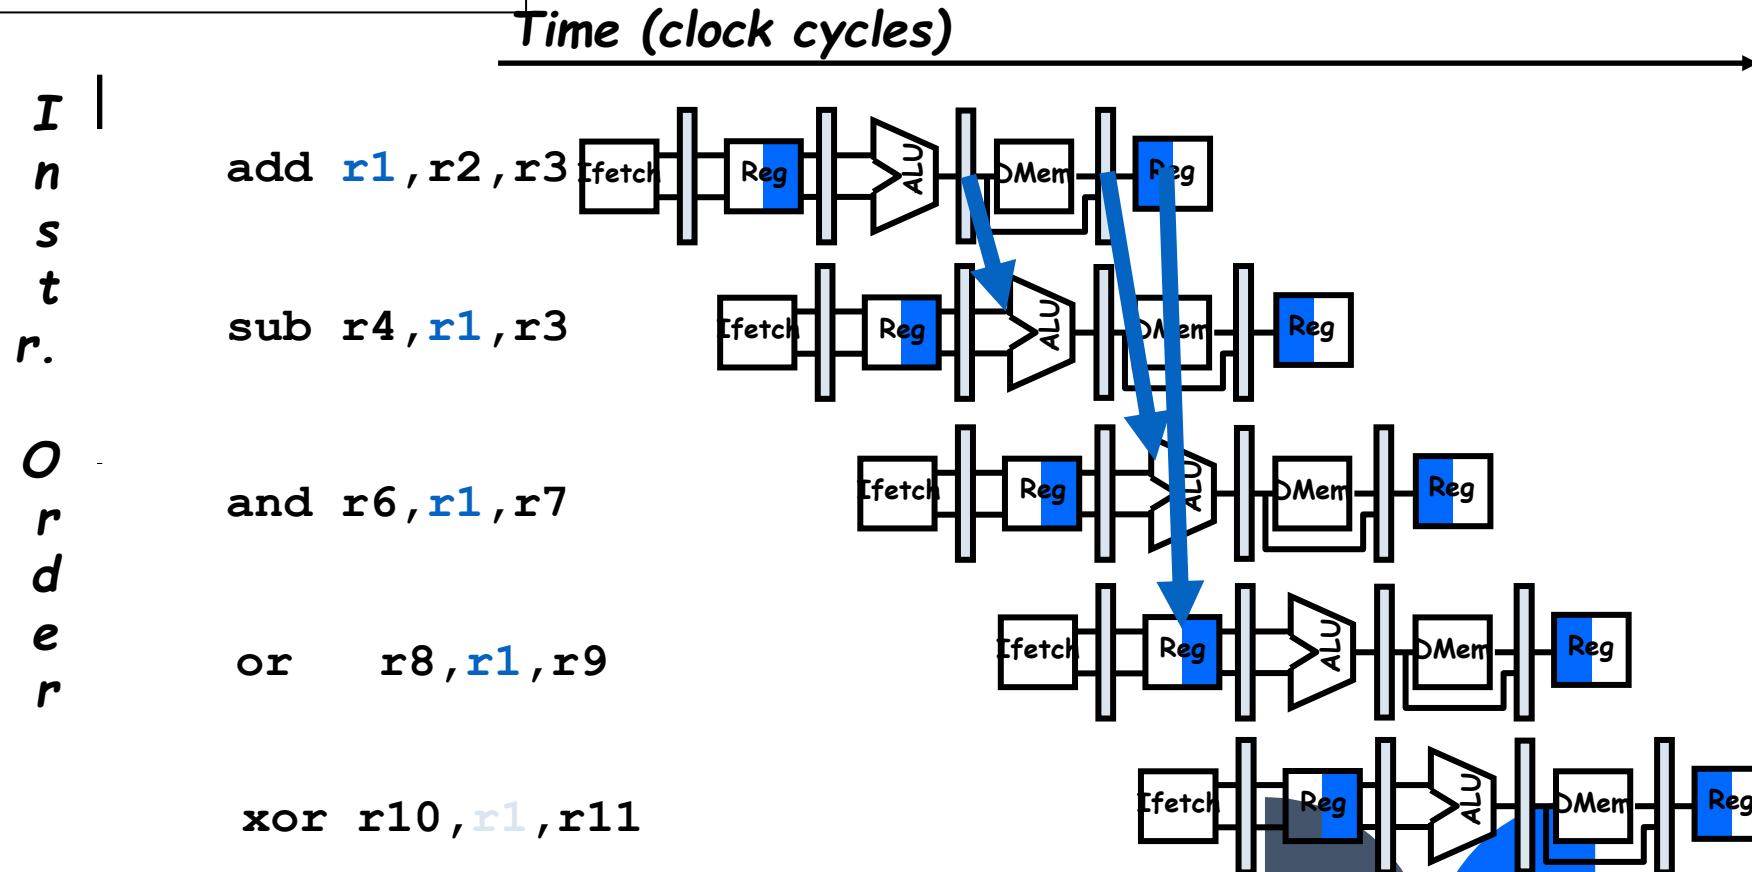


The use of the result of the ADD instruction in the next three instructions causes a hazard, since the register is not written until after those instructions read it.

Data Hazards

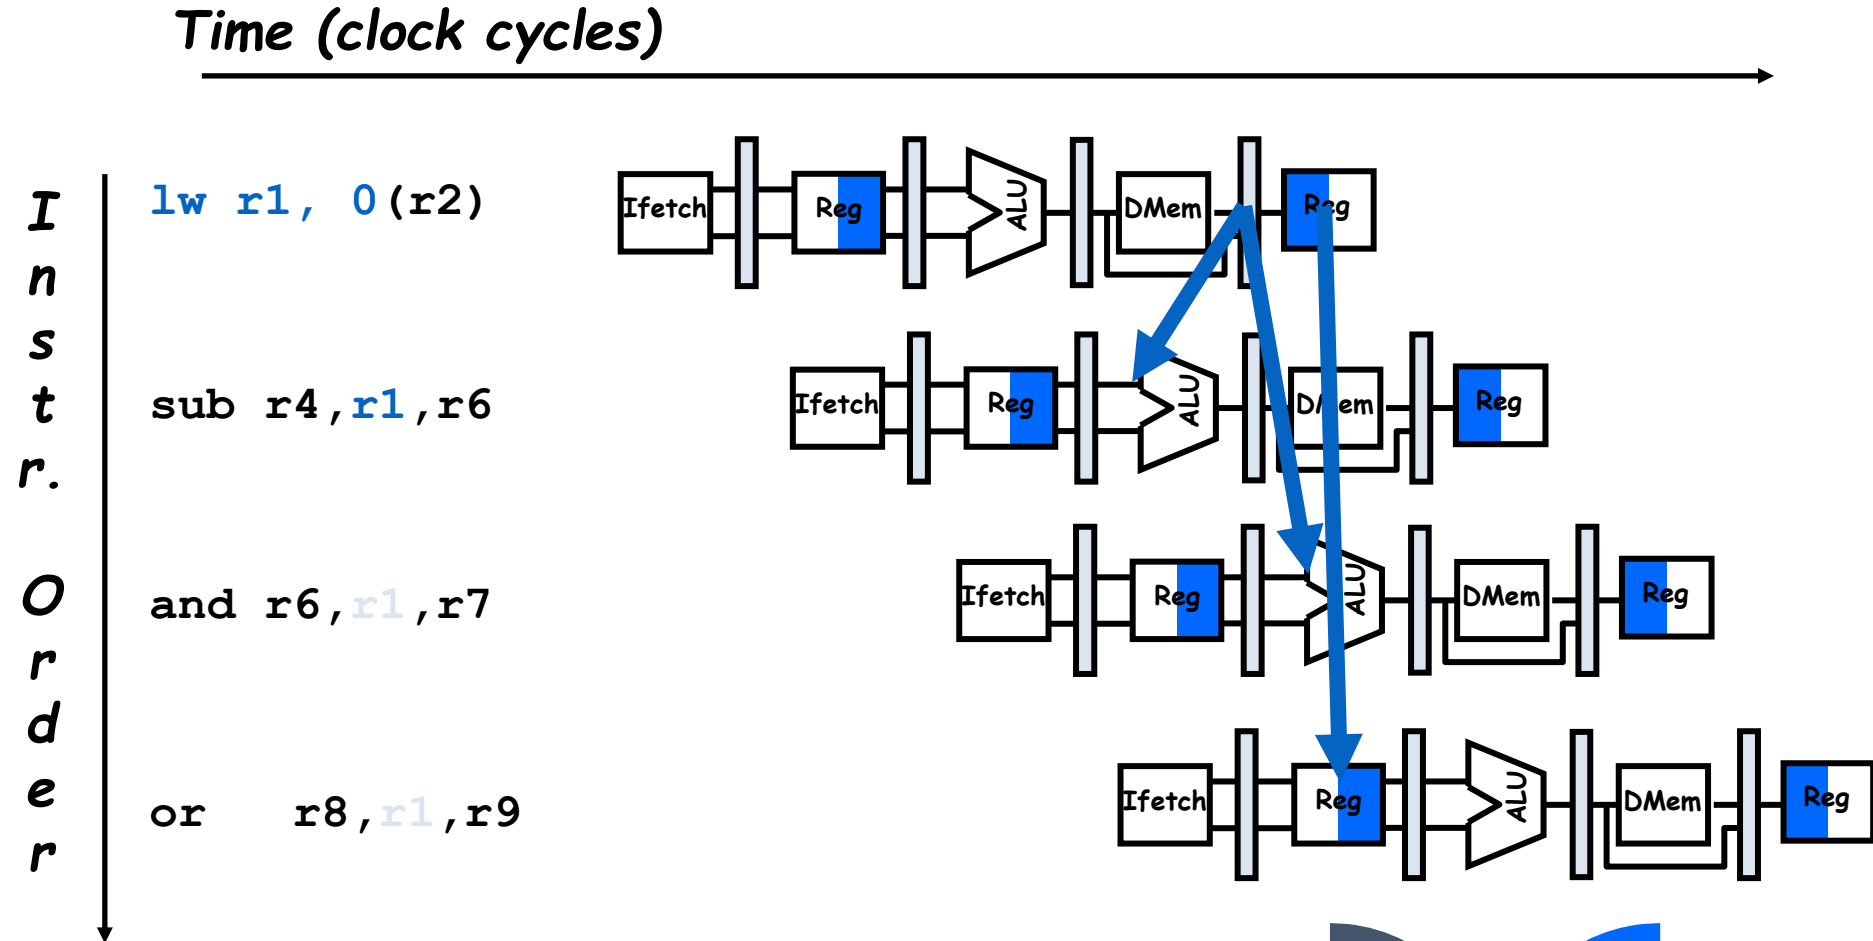
Forwarding To Avoid
Data Hazard

Forwarding is the concept of making data available to the input of the ALU for subsequent instructions, even though the generating instruction hasn't gotten to WB in order to write the memory or registers.



Data Hazards

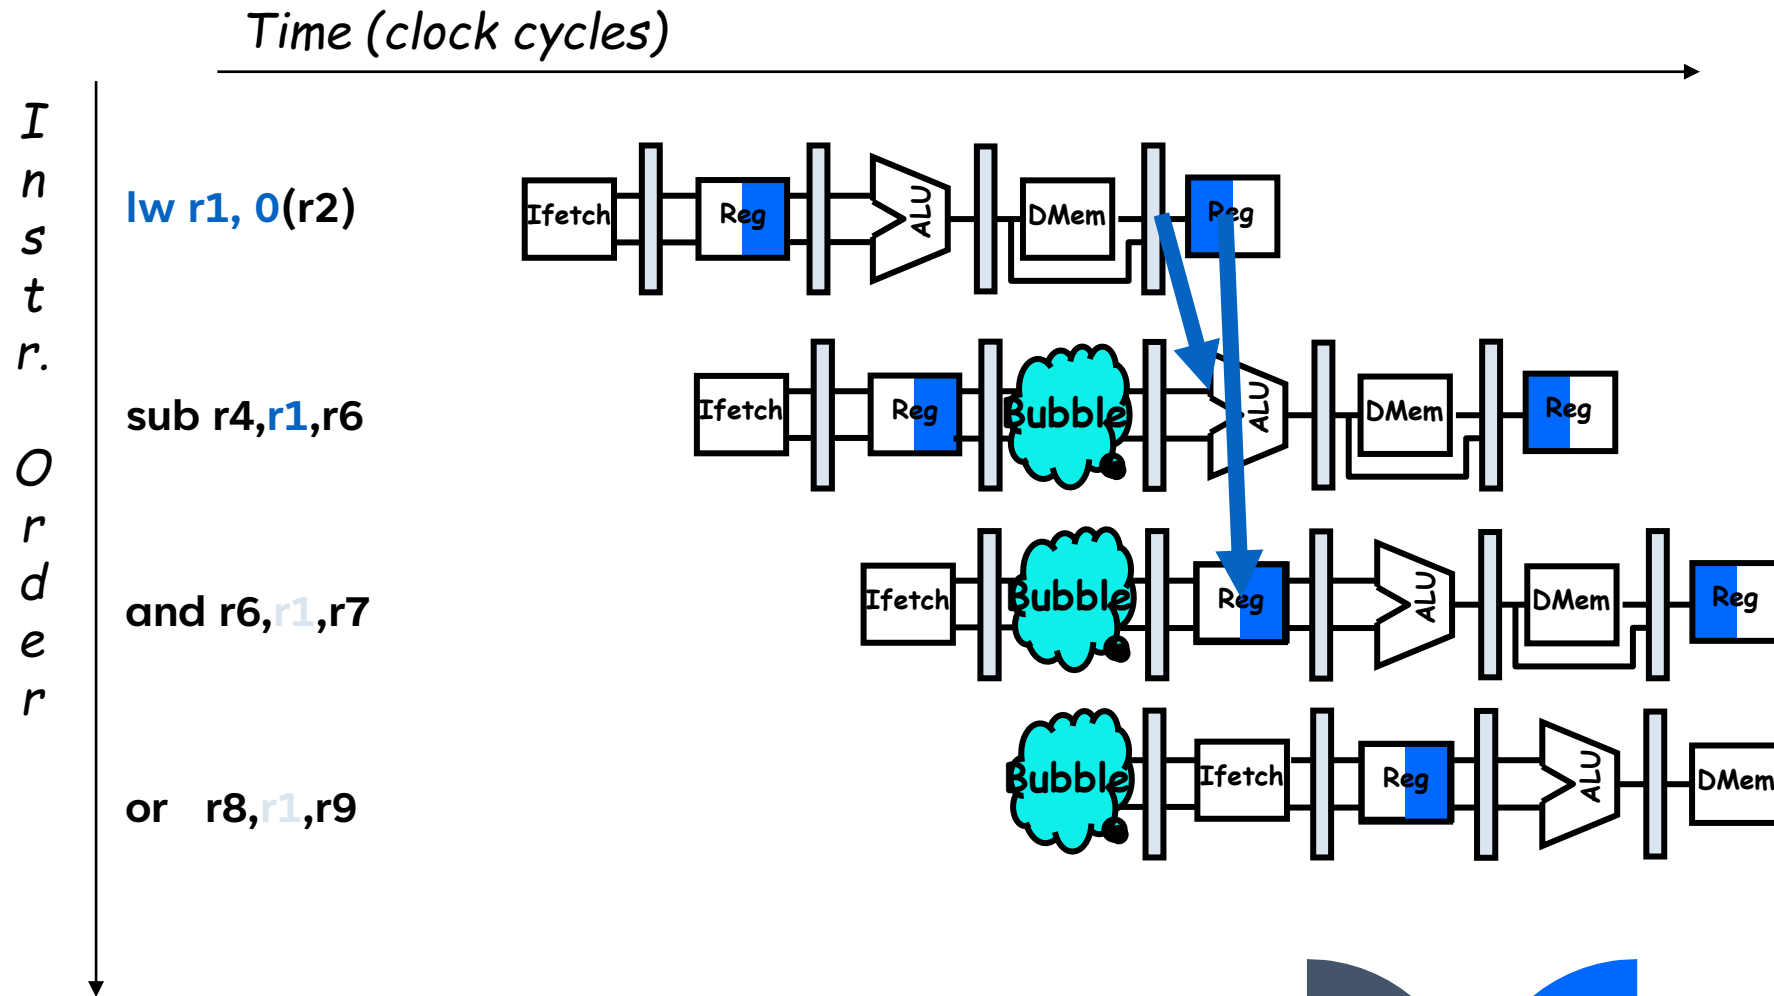
The data isn't loaded until after the MEM stage.



There are some instances where hazards occur, even with forwarding.

Data Hazards

The stall is necessary as shown here.



There are some instances where hazards occur, even with forwarding.

Data Hazards

This is another representation of the stall.

LW	R1, 0(R2)	IF	ID	EX	MEM	WB			
SUB	R4, R1, R5		IF	ID	EX	MEM	WB		
AND	R6, R1, R7			IF	ID	EX	MEM	WB	
OR	R8, R1, R9				IF	ID	EX	MEM	WB

LW	R1, 0(R2)	IF	ID	EX	MEM	WB				
SUB	R4, R1, R5		IF	ID	stall	EX	MEM	WB		
AND	R6, R1, R7			IF	stall	ID	EX	MEM	WB	
OR	R8, R1, R9				stall	IF	ID	EX	MEM	WB

Data Hazards

Pipeline Scheduling

Instruction scheduled by compiler - move instruction in order to reduce stall.


lw Rb, b	code sequence for $a = b+c$ before scheduling
lw Rc, c	
Add Ra, Rb, Rc	stall
sw a, Ra	
lw Re, e	code sequence for $d = e+f$ before scheduling
lw Rf, f	
sub Rd, Re, Rf	stall
sw d, Rd	

Arrangement of code after scheduling.

```
lw Rb, b
lw Rc, c
lw Re, e
Add Ra, Rb, Rc
lw Rf, f
sw a, Ra
sub Rd, Re, Rf
sw d, Rd
```

Data Hazards

Pipeline Scheduling



Control Hazards

A.1 What is Pipelining?

A.2 The Major Hurdle of Pipelining- Structural Hazards

- Structural Hazards
- Data Hazards
- Control Hazards

A.3 How is Pipelining Implemented


A.4 What Makes Pipelining Hard to Implement?

A.5 Extending the MIPS Pipeline to Handle Multi-cycle Operations

A control hazard is when we need to find the destination of a branch, and can't fetch any new instructions until we know that destination.

Control Hazards

Control Hazard on Branches Three Stage Stall



Control Hazards

Branch Stall Impact

If CPI = 1, 30% branch, Stall 3 cycles => new CPI = 1.9!

(Whoa! How did we get that 1.9???)

Two part solution to this dramatic increase:

- Determine branch taken or not sooner, AND

- Compute taken branch address earlier

MIPS branch tests if register = 0 or $\neq 0$

MIPS Solution:

- Move Zero test to ID/RF stage

- Adder to calculate new PC in ID/RF stage

 - must be fast

 - can't afford to subtract

 - compares with 0 are simple

 - Greater-than, Less-than test signbit, but not-equal must OR all bits

 - more general compares need ALU

- 1 clock cycle penalty for branch versus 3

In the next chapter, we'll look at ways to avoid the branch all together.

Control Hazards

Five Branch Hazard Alternatives

#1: Stall until branch direction is clear

#2: Predict Branch Not Taken

Execute successor instructions in sequence

“Squash” instructions in pipeline if branch actually taken

Advantage of late pipeline state update

47% MIPS branches not taken on average

PC+4 already calculated, so use it to get next instruction

#3: Predict Branch Taken

53% MIPS branches taken on average

But haven't calculated branch target address in MIPS

MIPS still incurs 1 cycle branch penalty

Other machines: branch target known before outcome

Control Hazards

Five Branch Hazard Alternatives

#4: Execute Both Paths

#5: Delayed Branch

Define branch to take place **AFTER** a following instruction

```
branch instruction
  sequential successor1
  sequential successor2
  .....
  sequential successorn
branch target if taken
```

Branch delay of length n

1 slot delay allows proper decision and branch target address in 5 stage pipeline

MIPS uses this

Control Hazards

Delayed Branch

Where to get instructions to fill branch delay slot?

- Before branch instruction

- From the target address: only valuable when branch taken

- From fall through: only valuable when branch not taken

- Cancelling branches allow more slots to be filled

Compiler effectiveness for single branch delay slot:

- Fills about 60% of branch delay slots

- About 80% of instructions executed in branch delay slots useful in computation

- About 50% ($60\% \times 80\%$) of slots usefully filled

Delayed Branch downside: 7-8 stage pipelines, multiple instructions issued per clock (superscalar)

Control Hazards

Evaluating Branch Alternatives

$$\text{Pipeline speedup} = \frac{\text{Pipeline depth}}{1 + \text{Branch frequency} \times \text{Branch penalty}}$$

<i>Scheduling scheme</i>	<i>Branch penalty</i>	<i>CPI</i>	<i>speedup v. unpipelined</i>	<i>Speedup v. stall</i>
Stall pipeline	3	1.42	3.5	1.0
Predict taken	1	1.14	4.4	1.26
Predict not taken	1	1.09	4.5	1.29
Delayed branch	0.5	1.07	4.6	1.31

Conditional & Unconditional = 14%,

65% change PC

Control Hazards

Pipelining Introduction Summary

Just overlap tasks, and easy if tasks are independent

Speed Up \propto Pipeline Depth; if ideal CPI is 1, then:

$$\text{Speedup} = \frac{\text{Pipeline Depth}}{1 + \text{Pipeline stall CPI}} \times \frac{\text{Clock Cycle Unpipelined}}{\text{Clock Cycle Pipelined}}$$

Hazards limit performance on computers:

Structural: need more HW resources


Data (RAW,WAR,WAW): need forwarding, compiler scheduling

Control: delayed branch, prediction


Control Hazards

The compiler can program what it thinks the branch direction will be. Here are the results when it does so.

Compiler “Static”
Prediction of
Taken/Untaken Branches



Always taken



Taken backwards
Not Taken Forwards

Control Hazards

Compiler “Static” Prediction of Taken/Untaken Branches

Improves strategy for placing instructions in delay slot

Two strategies

- Backward branch predict taken, forward branch not taken


- Profile-based prediction: record branch behavior, predict branch based on prior run

Control Hazards

Evaluating Static Branch Prediction Strategies

Misprediction ignores frequency of branch

“Instructions between mispredicted branches” is a better metric



What Makes Pipelining Hard?

A.1 What is Pipelining?

**A.2 The Major Hurdle of Pipelining-
Structural Hazards**

- Data Hazards
- Control Hazards

A.3 How is Pipelining Implemented

**A.4 What Makes Pipelining Hard to
Implement?**

**A.5 Extending the MIPS Pipeline to
Handle Multi-cycle Operations**

What Makes Pipelining Hard?

Examples of interrupts:

Power failing,
Arithmetic overflow,
I/O device request,
OS call,
Page fault

Interrupts (also known as: faults,
exceptions, traps) often require

surprise jump (to vectored address)
linking return address
saving of PSW (including CCs)
state change (e.g., to kernel mode)

Interrupts cause
great havoc!

**There are 5 instructions executing
in 5 stage pipeline when an
interrupt occurs:**

- **How to stop the pipeline?**
- **How to restart the pipeline?**
- **Who caused the interrupt?**

What Makes Pipelining Hard?

Interrupts cause
great havoc!

What happens on interrupt while in delay slot ?

- Next instruction is not sequential
- solution #1: save multiple PCs
- Save current and next PC
 - Special return sequence, more complex hardware
- solution #2: single PC plus
- Branch delay bit
 - PC points to branch instruction

Stage

Problem that causes the interrupt

IF

Page fault on instruction fetch; misaligned memory access; memory-protection violation

ID

Undefined or illegal opcode

EX

Arithmetic interrupt

MEM

Page fault on data fetch; misaligned memory access; memory-protection violation

What Makes Pipelining Hard?

Interrupts cause
great havoc!

Simultaneous exceptions in more than one pipeline stage, e.g.,

- Load with data page fault in MEM stage

- Add with instruction page fault in IF stage

- Add fault will happen BEFORE load fault

Solution #1

- Interrupt status vector per instruction

- Defer check until last stage, kill state update if exception

Solution #2

- Interrupt ASAP

- Restart everything that is incomplete

Another advantage for state update late in pipeline!

What Makes Pipelining Hard?

Interrupts cause
great havoc!

Here's what happens on a data page fault.

	1	2	3	4	5	6	7	8	9
i	F	D	X	M	W				
i+1		F	D	X	M	W	< page fault		
i+2			F	D	X	M	W	< squash	
i+3				F	D	X	M	W	< squash
i+4					F	D	X	M	W < squash
i+5	trap >					F	D	X	M W
i+6	trap handler >						F	D	X M W

What Makes Pipelining Hard?

Complex Instructions

Complex Addressing Modes and Instructions

Address modes: Autoincrement causes register change during instruction execution

Interrupts? Need to restore register state

Adds WAR and WAW hazards since writes are no longer the last stage.

Memory-Memory Move Instructions

Must be able to handle multiple page faults

Long-lived instructions: partial state save on interrupt

Condition Codes

Handling Multi-cycle Operations

A.1 What is Pipelining?

A.2 The Major Hurdle of Pipelining- Structural Hazards

- Data Hazards
- Control Hazards

A.3 How is Pipelining Implemented

A.4 What Makes Pipelining Hard to Implement?

A.5 Extending the MIPS Pipeline to Handle Multi-cycle Operations

Multi-cycle instructions also lead to pipeline complexity.

A very lengthy instruction causes everything else in the pipeline to wait for it.

Multi-Cycle Operations

Floating Point

Floating point gives long execution time.

This causes a stall of the pipeline.

It's possible to pipeline the FP execution unit so it can initiate new instructions without waiting full latency. Can also have multiple FP units.

<i>FP Instruction</i>	<i>Latency</i>	<i>Initiation Rate</i>
Add, Subtract	4	3
Multiply	8	4
Divide	36	35
Square root	112	111
Negate	2	1
Absolute value	2	1
FP compare	3	2

Multi-Cycle Operations

Floating Point

Divide, Square Root take 10X to 30X longer than Add

Interrupts?

Adds WAR and WAW hazards since pipelines are no longer same length

	1	2	3	4	5	6	7	8	9	10	11
i	IF	ID	EX	MEM	WB						
i + 1		IF	ID	EX	EX	EX	EX	MEM	WB		
i + 2			IF	ID	EX	MEM	WB				
i + 3				IF	ID	EX	EX	EX	EX	MEM	WB
i + 4					IF	ID	EX	MEM	WB		
i + 5						IF	ID	--	--	EX	EX
i + 6							IF	--	--	ID	EX

Notes:

i + 2: no WAW, but this complicates an interrupt

i + 4: no WB conflict

i + 5: stall forced by structural hazard

i + 6: stall forced by in-order issue

Summary of Pipelining Basics

Hazards limit performance

- Structural: need more HW resources

- Data: need forwarding, compiler scheduling

- Control: early evaluation & PC, delayed branch, prediction

Increasing length of pipe increases impact of hazards; pipelining helps instruction bandwidth, not latency

Interrupts, Instruction Set, FP makes pipelining harder

Compilers reduce cost of data and control hazards

- Load delay slots

- Branch delay slots

- Branch prediction

Summary

A.1 What is Pipelining?

A.2 The Major Hurdle of Pipelining-Structural Hazards

- Data Hazards
- Control Hazards

A.3 How is Pipelining Implemented

A.4 What Makes Pipelining Hard to Implement?

A.5 Extending the MIPS Pipeline to Handle Multi-cycle Operations