


MODULE 3: MEMORY



Nilanjan Byabarta
University of Engineering and Management, Kolkata

OBJECTIVES

- Master the concepts of hierarchical memory organization.
- Understand how each level of memory contributes to system performance, and how the performance is measured.
- Master the concepts behind cache memory, virtual memory, memory segmentation, paging and address translation.

INTRODUCTION

- Memory lies at the heart of the stored-program computer.

TYPES OF MEMORY

- There are two kinds of main memory: *random access memory, RAM, and read-only-memory, ROM.*
- There are two types of RAM, dynamic RAM (DRAM) and static RAM (SRAM).
- DRAM consists of capacitors that slowly leak their charge over time. Thus, they must be refreshed every few milliseconds to prevent data loss.
- DRAM is “cheap” memory owing to its simple design.

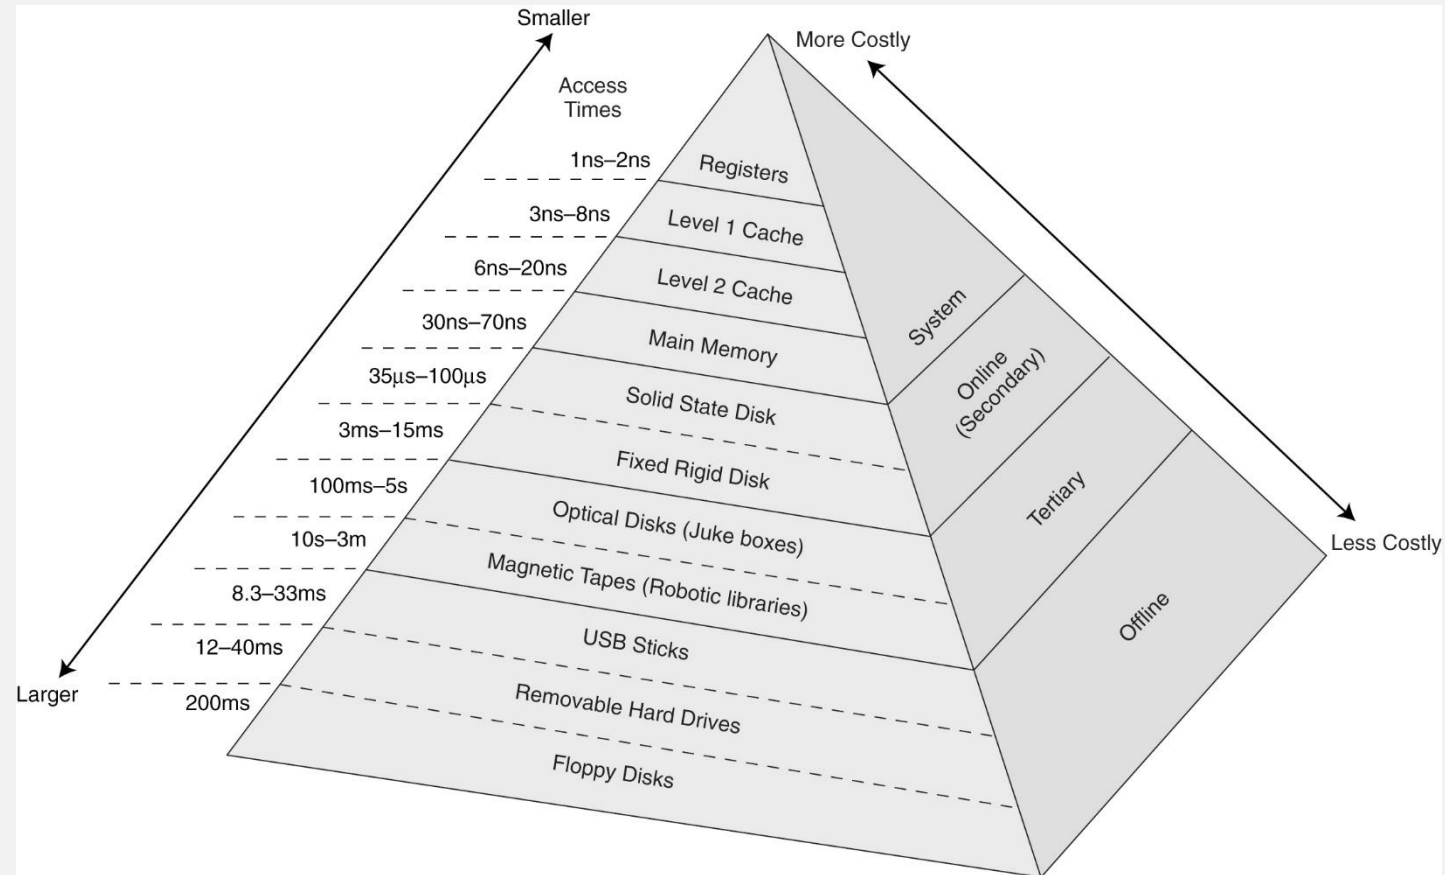
TYPES OF MEMORY

- SRAM consists of circuits similar to the D flip-flop
- SRAM is very fast memory and it doesn't need to be refreshed like DRAM does. It is used to build cache memory,
- ROM also does not need to be refreshed, either. In fact, it needs very little charge to retain its memory.
- ROM is used to store permanent, or semi-permanent data that persists even while the system is turned off.

THE MEMORY HIERARCHY

- Generally speaking, faster memory is more expensive than slower memory.
- To provide the best performance at the lowest cost, memory is organized in a hierarchical fashion.
- Small, fast storage elements are kept in the CPU, larger, slower main memory is accessed through the data bus.
- Larger, (almost) permanent storage in the form of disk and tape drives is still further from the CPU.

THE MEMORY HIERARCHY



- This storage organization can be thought of as a pyramid:

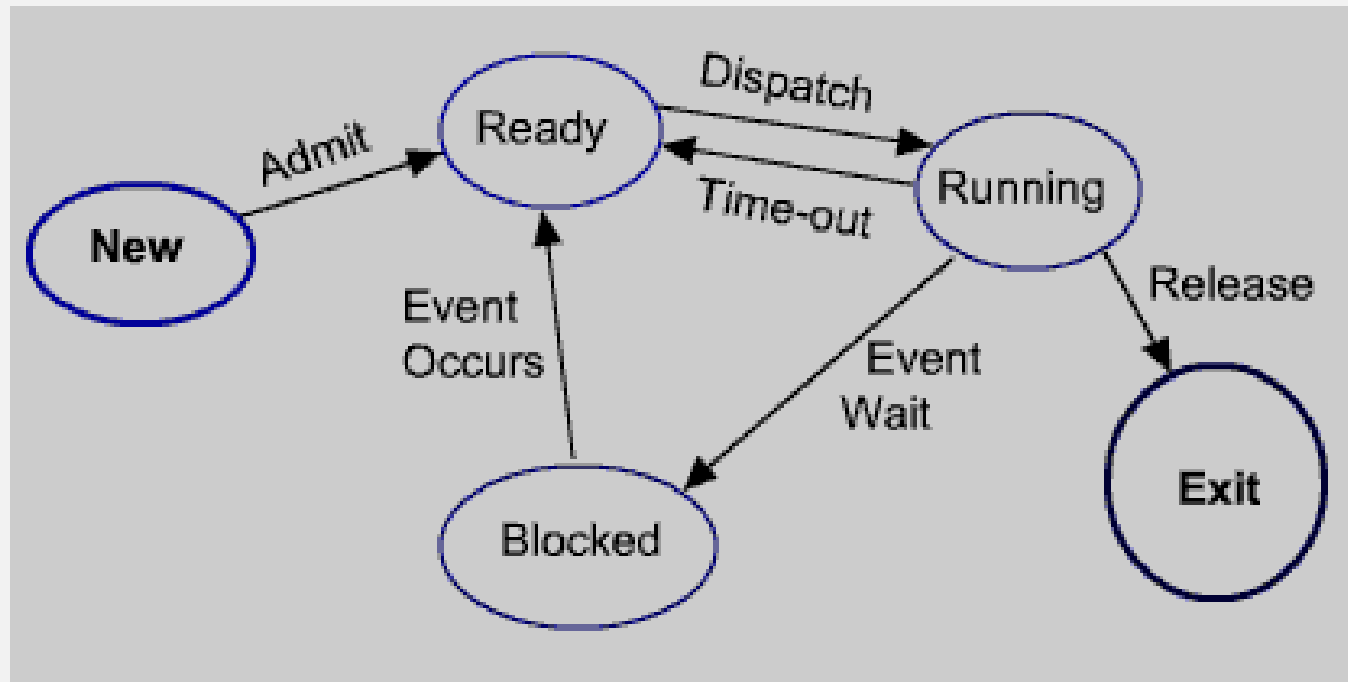
THE MEMORY HIERARCHY

- We are most interested in the memory hierarchy that involves registers, cache, main memory, and virtual memory.
- Registers are storage locations available on the processor itself.
- Virtual memory is typically implemented using a hard drive; it extends the address space from RAM to the hard drive.
- Virtual memory provides more space: Cache memory provides speed.

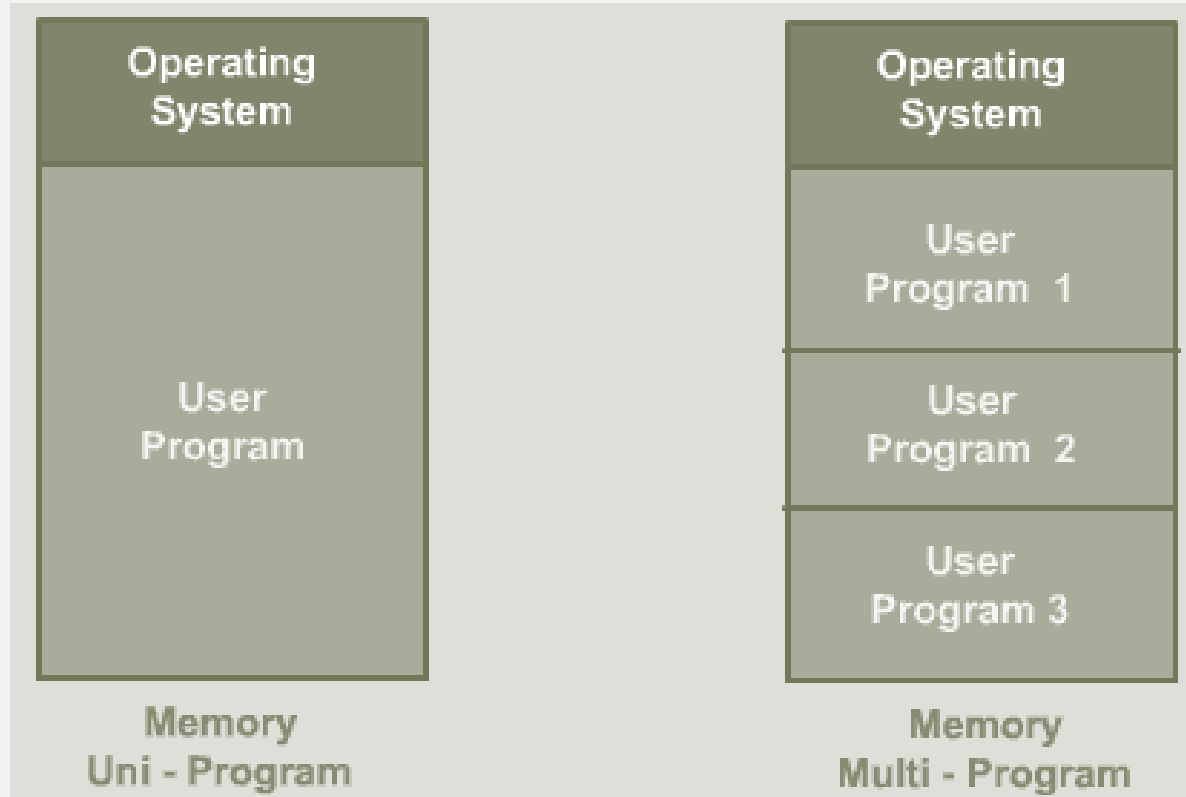
THE MEMORY HIERARCHY

- To access a particular piece of data, the CPU first sends a request to its nearest memory, usually cache.
- If the data is not in cache, then main memory is queried. If the data is not in main memory, then the request goes to disk.
- Once the data is located, then the data, and a number of its nearby data elements are fetched into cache memory.

MEMORY CYCLE



MEMORY MANAGEMENT



MEMORY PARTITIONS

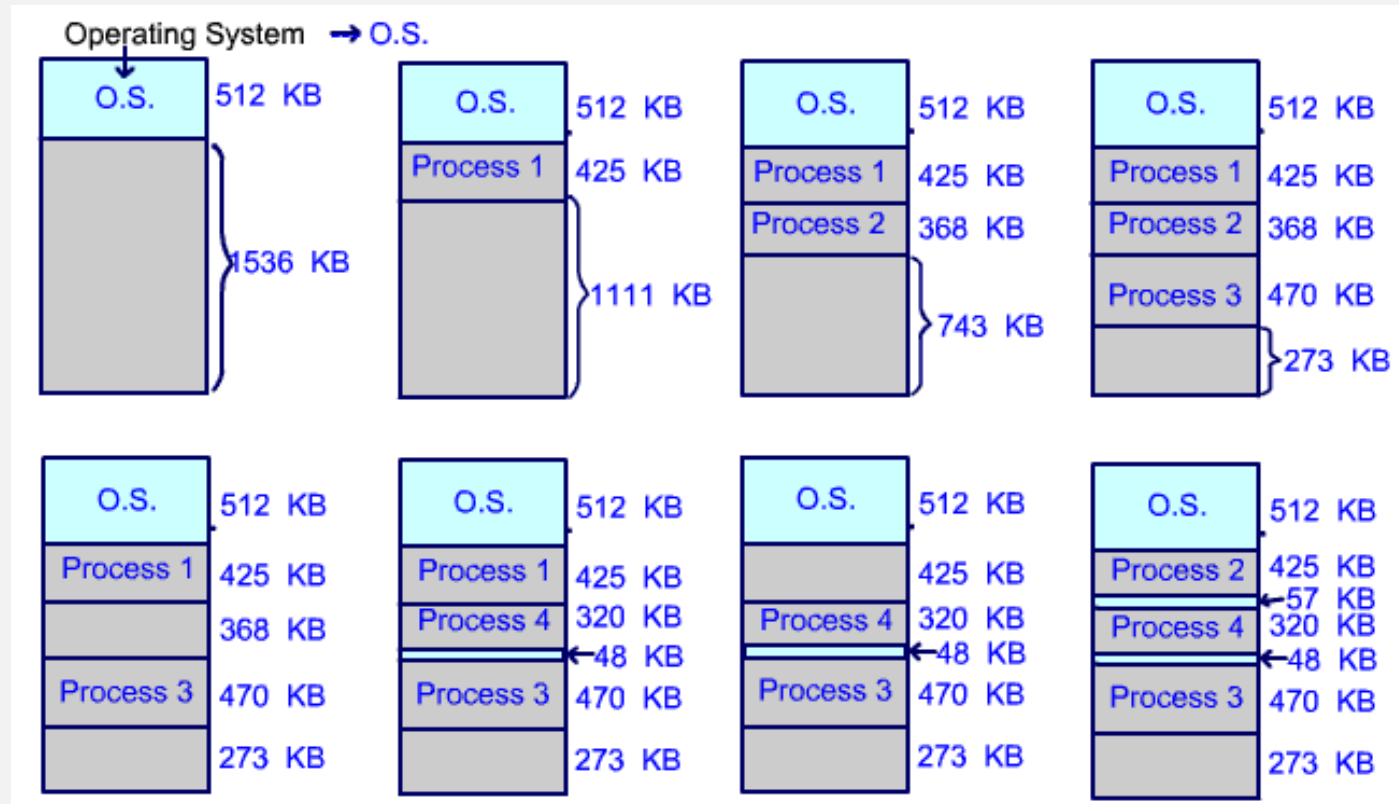
8 M Operating System
8 M
8 M
8 M
8 M
8 M
8 M
8 M

(a) Equal Size Partitions

8 M Operating System
6 M
16 M
13 M
5 M
8 M
8 M

(b) Un-equal Size Partitions

MEMORY DYNAMIC PARTITIONS



THE MEMORY HIERARCHY

- This leads us to some definitions.
 - A *hit* is when data is found at a given memory level.
 - A *miss* is when it is not found.
 - The *hit rate* is the percentage of time data is found at a given memory level.
 - The *miss rate* is the percentage of time it is not.
 - Miss rate = $1 - \text{hit rate}$.
 - The *hit time* is the time required to access data at a given memory level.
 - The *miss penalty* is the time required to process a miss, including the time that it takes to replace a block of memory plus the time it takes to deliver the data to the processor.

THE MEMORY HIERARCHY


- An entire blocks of data is copied after a hit because the *principle of locality* tells us that once a byte is accessed, it is likely that a nearby data element will be needed soon.
- There are three forms of locality:
 - *Temporal locality*- Recently-accessed data elements tend to be accessed again.
 - *Spatial locality* - Accesses tend to cluster.
 - *Sequential locality* - Instructions tend to be accessed sequentially.

CACHE MEMORY

- The purpose of cache memory is to speed up accesses by storing recently used data closer to the CPU, instead of storing it in main memory.
- Although cache is much smaller than main memory, its access time is a fraction of that of main memory.
- Unlike main memory, which is accessed by address, cache is typically accessed by content; hence, it is often called *content addressable memory*.
- Because of this, a single large cache memory isn't always desirable-- it takes longer to search.

CACHE MEMORY

- The “content” that is addressed in content addressable cache memory is a subset of the bits of a main memory address called a *field*.
- Many blocks of main memory map to a single block of cache. A *tag* field in the cache block distinguishes one cached memory block from another.
- A *valid bit* indicates whether the cache block is being used.
- An *offset field* points to the desired data in the block.




CACHE MEMORY

- The simplest cache mapping scheme is *direct mapped cache(Direct Mapping)*.
- In a direct mapped cache consisting of N blocks of cache, block X of main memory maps to cache block $Y = X \bmod N$.
- Thus, if we have 10 blocks of cache, block 7 of cache may hold blocks 7, 17, 27, 37, . . . of main memory.

CACHE MEMORY

- With direct mapped cache consisting of N blocks of cache, block X of main memory maps to cache block $Y = X \bmod N$.




CACHE MEMORY

- EXAMPLE 1: Consider a word-addressable main memory consisting of four blocks, and a cache with two blocks, where each block is 4 words.
- This means Block 0 and 2 of main memory map to Block 0 of cache, and Blocks 1 and 3 of main memory map to Block 1 of cache.
- Using the tag, block, and offset fields, we can see how main memory maps to cache as follows.

CACHE MEMORY

- EXAMPLE 1 Consider a word-addressable main memory consisting of four blocks, and a cache with two blocks, where each block is 4 words.
- First, we need to determine the address format for mapping. Each block is 4 words, so the offset field must contain 2 bits; there are 2 blocks in cache, so the block field must contain 1 bit; this leaves 1 bit for the tag (as a main memory address has 4 bits because there are a total of $2^4=16$ words).




CACHE MEMORY


- EXAMPLE 1 Cont'd

- Suppose we need to access main memory address 3_{16} (0011 in binary). If we partition 0011 using the address format from Figure a, we get Figure b.
- Thus, the main memory address 0011 maps to cache block 0.
- Figure c shows this mapping, along with the tag that is also stored with the data.


a







b



c



CACHE MEMORY

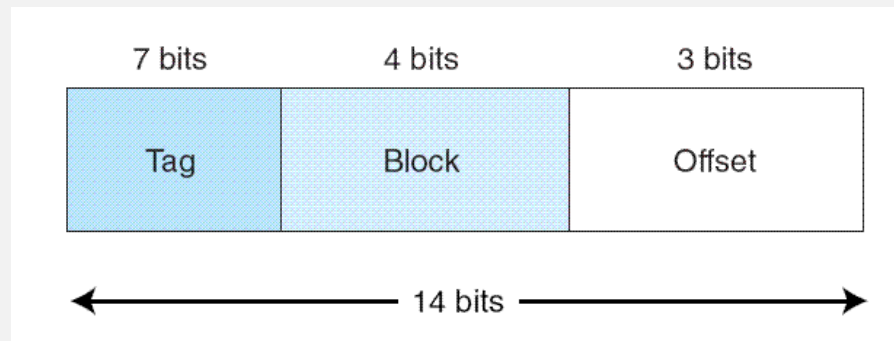


CACHE MEMORY

- EXAMPLE 2 Assume a byte-addressable memory consists of 2^{14} bytes, cache has 16 blocks, and each block has 8 bytes.

$$\frac{2^{14}}{2^3} = 2^{11}$$

- The number of memory blocks are:
- Each main memory address requires 14 bits. Of this 14-bit address field, the rightmost 3 bits reflect the offset field
- We need 4 bits to select a specific block in cache, so the block field consists of the middle 4 bits.
- The remaining 7 bits make up the tag field.



CACHE MEMORY


- In summary, direct mapped cache maps main memory blocks in a modular fashion to cache blocks. The mapping depends on:
- The number of bits in the main memory address (how many addresses exist in main memory)
- The number of blocks are in cache (which determines the size of the block field)
- How many addresses (either bytes or words) are in a block (which determines the size of the offset field)

CACHE MEMORY

- Suppose instead of placing memory blocks in specific cache locations based on memory address, we could allow a block to go anywhere in cache.
- In this way, cache would have to fill up before any blocks are evicted.
- This is how *fully associative cache Mapping* works.
- A memory address is partitioned into only two fields: the tag and the word.

CACHE MEMORY

- Suppose, as before, we have 14-bit memory addresses and a cache with **16 blocks**, each **block of size 8**. The field format of a memory reference is:



- When the cache is searched, all tags are searched in parallel to retrieve the data quickly.
- This requires special, costly hardware.

CACHE MEMORY

- You will recall that direct mapped cache evicts a block whenever another memory reference needs that block.
- With fully associative cache, we have no such mapping, thus we must devise an algorithm to determine which block to evict from the cache.
- The block that is evicted is the *victim block*.
- There are a number of ways to pick a *victim*

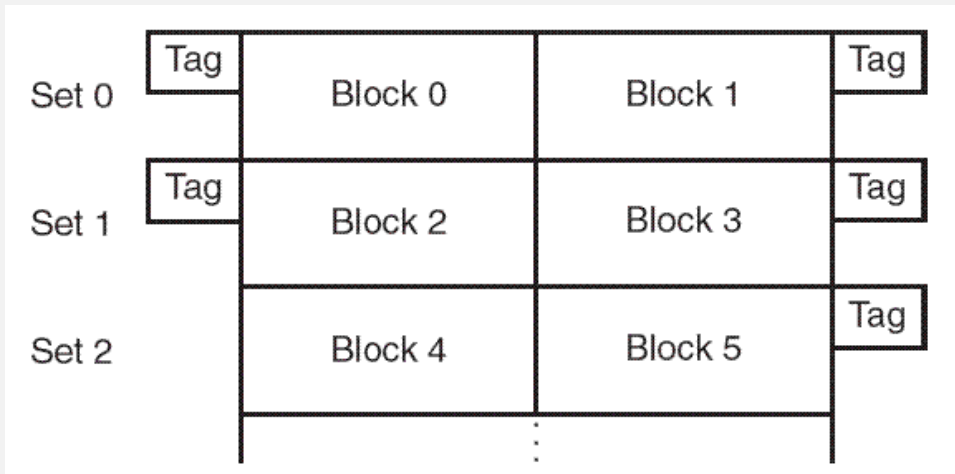
CACHE MEMORY

- **Set associative cache Mapping** combines the ideas of direct mapped cache and fully associative cache.
- An N -way set associative cache mapping is like direct mapped cache in that a memory reference maps to a particular location in cache.
- Unlike direct mapped cache, a memory reference maps to a set of several cache blocks, similar to the way in which fully associative cache works.
- Instead of mapping anywhere in the entire cache, a memory reference can map only to the subset of cache slots.


CACHE MEMORY

- The number of cache blocks per set in set associative cache varies according to overall system design.
 - For example, a 2-way set associative cache can be conceptualized as shown in the schematic below.
 - Each set contains two different memory blocks.

Logical view



Linear view



CACHE MEMORY

- In set associative cache mapping, a memory reference is divided into three fields: tag, set, and offset.
- As with direct-mapped cache, the offset field chooses the word within the cache block, and the tag field uniquely identifies the memory address.
- The set field determines the set to which the memory block maps.

CACHE MEMORY

- With fully associative and set associative cache, a *replacement policy* is invoked when it becomes necessary to evict a block from cache.
- An *optimal* replacement policy would be able to look into the future to see which blocks won't be needed for the longest period of time.
- Although it is impossible to implement an optimal replacement algorithm, it is instructive to use it as a benchmark for assessing the efficiency of any other scheme we come up with.

CACHE MEMORY

- The replacement policy that we choose depends upon the locality that we are trying to optimize-- usually, we are interested in **temporal locality**.
- A ***least recently used* (LRU) algorithm** keeps track of the last time that a block was accessed and evicts the block that has been unused for the longest period of time.
- The disadvantage of this approach is its complexity: **LRU has to maintain an access history for each block**, which ultimately slows down the cache.

CACHE MEMORY

- *First-in, first-out (FIFO)* is a popular cache replacement policy.
- In FIFO, the block that has been in the cache the longest, regardless of when it was last used.
- *A random replacement policy* does what its name implies: It picks a block at random and replaces it with a new block.
- Random replacement can certainly evict a block that will be needed often or needed soon, but it never **thrashes**.

CACHE MEMORY

- The performance of hierarchical memory is measured by its *effective access time (EAT)*.
- EAT is a weighted average that takes into account the hit ratio and relative access times of successive levels of memory.
- The EAT for a two-level memory is given by:

$$EAT = H \times Access_C + (1-H) \times Access_{MM}.$$

where H is the cache hit rate and $Access_C$ and $Access_{MM}$ are the access times for cache and main memory, respectively.

CACHE MEMORY

- For example, consider a system with a main memory access time of 200ns supported by a cache having a 10ns access time and a hit rate of 99%. (For non Overlapping)
- Suppose access to cache and main memory occurs concurrently. (The accesses overlap.)
- The EAT is:

$$0.99(10\text{ns}) + 0.01(200\text{ns}) = 9.9\text{ns} + 2\text{ns} = 11\text{ns}.$$

CACHE MEMORY

- For example, consider a system with a main memory access time of 200ns supported by a cache having a 10ns access time and a hit rate of 99%.
- If the accesses do not overlap, the EAT is:

$$\begin{aligned} &0.99(10\text{ns}) + 0.01(10\text{ns} + 200\text{ns}) \\ &= 9.9\text{ns} + 2.01\text{ns} = 12\text{ns}. \end{aligned}$$

- This equation for determining the effective access time can be extended to any number of memory levels, as we will see in later sections.

CACHE MEMORY

- Caching is depends upon programs exhibiting good locality.
 - Some object-oriented programs have poor locality owing to their complex, dynamic structures.
 - Arrays stored in column-major rather than row-major order can be problematic for certain cache organizations.
- With poor locality, caching can actually cause performance degradation rather than performance improvement.

CACHE MEMORY

- Cache replacement policies must take into account *dirty blocks*, those blocks that have been updated while they were in the cache.
- Dirty blocks must be written back to memory. *A write policy determines how this will be done.*
- There are two types of *write policies*, *write through* and *write back*.
- *Write through* updates cache and main memory simultaneously on every write.

CACHE MEMORY

- Write back (also called *copyback*) updates memory only when the block is selected for replacement.
- The disadvantage of write through is that memory must be updated with each cache write, which slows down the access time on updates. This slowdown is usually negligible, because the majority of accesses tend to be reads, not writes.
- The advantage of write back is that memory traffic is minimized, but its disadvantage is that memory does not always agree with the value in cache, causing problems in systems with many concurrent users.

CACHE MEMORY

- The cache we have been discussing is called a *unified* or *integrated* cache where both instructions and data are cached.
- Many modern systems employ **separate caches for data and instructions.**
 - **This is called a *Harvard* cache.**
- The separation of data from instructions provides better locality, at the cost of greater complexity.
 - Simply making the cache larger provides about the same performance improvement without the complexity.

CACHE MEMORY

- Cache performance can also be improved by adding a small associative cache to **hold blocks that have been evicted recently**.
 - This is called a *victim cache*.
- A *trace cache* is a variant of an instruction cache that holds decoded instructions for program branches, giving the illusion that noncontiguous instructions are really contiguous.

CACHE MEMORY

- Most of today's small systems employ multilevel cache hierarchies.
- The levels of cache form their own small memory hierarchy.
- Level1 cache (8KB to 64KB) is situated on the processor itself.
 - Access time is typically about 4ns.
- Level 2 cache (64KB to 2MB) may be on the motherboard, or on an expansion card.
 - Access time is usually around 15 - 20ns.

CACHE MEMORY

- In systems that employ three levels of cache, the Level 2 cache is placed on the same die as the CPU (reducing access time to about 10ns)
- Accordingly, the Level 3 cache (2MB to 256MB) refers to cache that is situated between the processor and main memory.
- Once the number of cache levels is determined, the next thing to consider is whether data (or instructions) can exist in more than one cache level.

CACHE MEMORY

- If the cache system used an *inclusive cache*, the **same data may be present at multiple levels of cache**.
- *Strictly inclusive caches* guarantee that **all data** in a smaller cache **also exists at the next higher level**.
- *Exclusive* caches permit **only one copy** of the data.
- The tradeoffs in choosing one over the other involve weighing the variables of access time, memory size, and circuit complexity.

VIRTUAL MEMORY

- Cache memory enhances performance by providing faster memory access speed.
- **Virtual memory** enhances performance by **providing greater memory capacity, without the expense of adding main memory.**
- Instead, a portion of a disk drive serves as an extension of main memory.
- If a system uses paging, virtual memory partitions main memory into individually managed *page frames*, that are written (*or paged*) to disk when they are not immediately needed.

VIRTUAL MEMORY


- A *physical address* is the *actual memory* address of physical memory.
- Programs create *virtual addresses* that are *mapped* to physical addresses by the memory manager.
- *Page faults* occur when a logical address requires that a page be brought in from disk.
- *Memory fragmentation* occurs when the paging process results in the creation of small, unusable clusters of memory addresses.

VIRTUAL MEMORY

- **Main memory and virtual memory** are divided into equal sized pages.
- The entire address space required by a process need not be in memory at once. Some parts can be on disk, while others are in main memory.
- Further, the pages allocated to a process do not need to be stored contiguously-- either on disk or in memory.
- In this way, only the needed pages are in memory at any time, the unnecessary pages are in slower disk storage.

VIRTUAL MEMORY

- Information concerning the location of each page, whether on disk or in memory, is maintained in a data structure called a *page table*.
- There is one page table for each active process.



VIRTUAL MEMORY


- When a process generates a virtual address, the operating system translates it into a physical memory address.
- To accomplish this, **the virtual address is divided into two fields: A *page* field, and an *offset* field.**
- The page field determines the page location of the address, and the offset indicates the location of the address within the page.
- The logical page number is translated into a physical page frame through a lookup in the page table.

VIRTUAL MEMORY

- If the valid bit is zero in the page table entry for the logical address, this means that the page is not in memory and must be fetched from disk.
 - This is a page fault.
 - If necessary, a page is evicted from memory and is replaced by the page retrieved from disk, and the valid bit is set to 1.
- If the valid bit is 1, the virtual page number is replaced by the physical frame number.
- The data is then accessed by adding the offset to the physical frame number.

VIRTUAL MEMORY

- As an example, suppose a system has a virtual address space of 8K and a physical address space of 4K, and the system uses byte addressing.
 - We have $2^{13}/2^{10} = 2^3$ virtual pages.
- A virtual address has 13 bits ($8K = 2^{13}$) with 3 bits for the page field and 10 for the offset, because the page size is 1024.
- A physical memory address requires 12 bits, the first two bits for the page frame and the trailing 10 bits the offset.



VIRTUAL MEMORY

- Suppose we have the page table shown below.
- What happens when CPU generates address 5459_{10}
 $= 1010101010011_2 = 1553_{16}$?

Page Table					
Page	Frame	Valid Bit	Addresses		
0	–	0	Page	Base 10	Base 16
1	3	1	0 :	0 – 1023	0 – 3FF
2	0	1	1 :	1024 – 2047	400 – 7FF
3	–	0	2 :	2048 – 3071	800 – BFF
4	–	0	3 :	3072 – 4095	C00 – FFF
5	1	1	4 :	4096 – 5119	1000 – 13FF
6	2	1	5 :	5120 – 6143	1400 – 17FF
7	–	0	6 :	6144 – 7167	1800 – 1BFF
			7 :	7168 – 8191	1C00 – 1FFF

VIRTUAL MEMORY

- We said earlier that **effective access time (EAT)** takes all levels of memory into consideration.
- Thus, virtual memory is also a factor in the calculation, and we also have to consider page table access time.
- Suppose a main memory access takes 200ns, the page fault rate is 1%, and it takes 10ms to load a page from disk. We have:

$$\text{EAT} = 0.99(200\text{ns} + 200\text{ns}) + 0.01(10\text{ms}) = 100,396\text{ns}.$$

VIRTUAL MEMORY

- Even if we had no page faults, the EAT would be 400ns because memory is always read twice: First to access the page table, and second to load the page from memory.
- Because page tables are read constantly, it makes sense to keep them in a special cache called a *translation look-aside buffer (TLB)*.
- TLBs are a special associative cache that stores the mapping of virtual pages to physical pages.

VIRTUAL MEMORY

- Another approach to virtual memory is the use of *segmentation*.
- Instead of dividing memory into equal-sized pages, virtual address space is divided into **variable-length segments**, often under the control of the programmer.
- A segment is located through its entry in a **segment table**, which contains the segment's memory location and a bounds limit that indicates its size.
- After a page fault, the operating system searches for a location in memory large enough to hold the segment that is retrieved from disk.

VIRTUAL MEMORY

- Both paging and segmentation can cause **fragmentation**.
- Paging is subject to *internal* fragmentation because a process may not need the entire range of addresses contained within the page. Thus, there may be many pages containing **unused fragments of memory**.
- Segmentation is subject to **external fragmentation**, which occurs when contiguous chunks of memory become broken up as segments are allocated and deallocated over time.

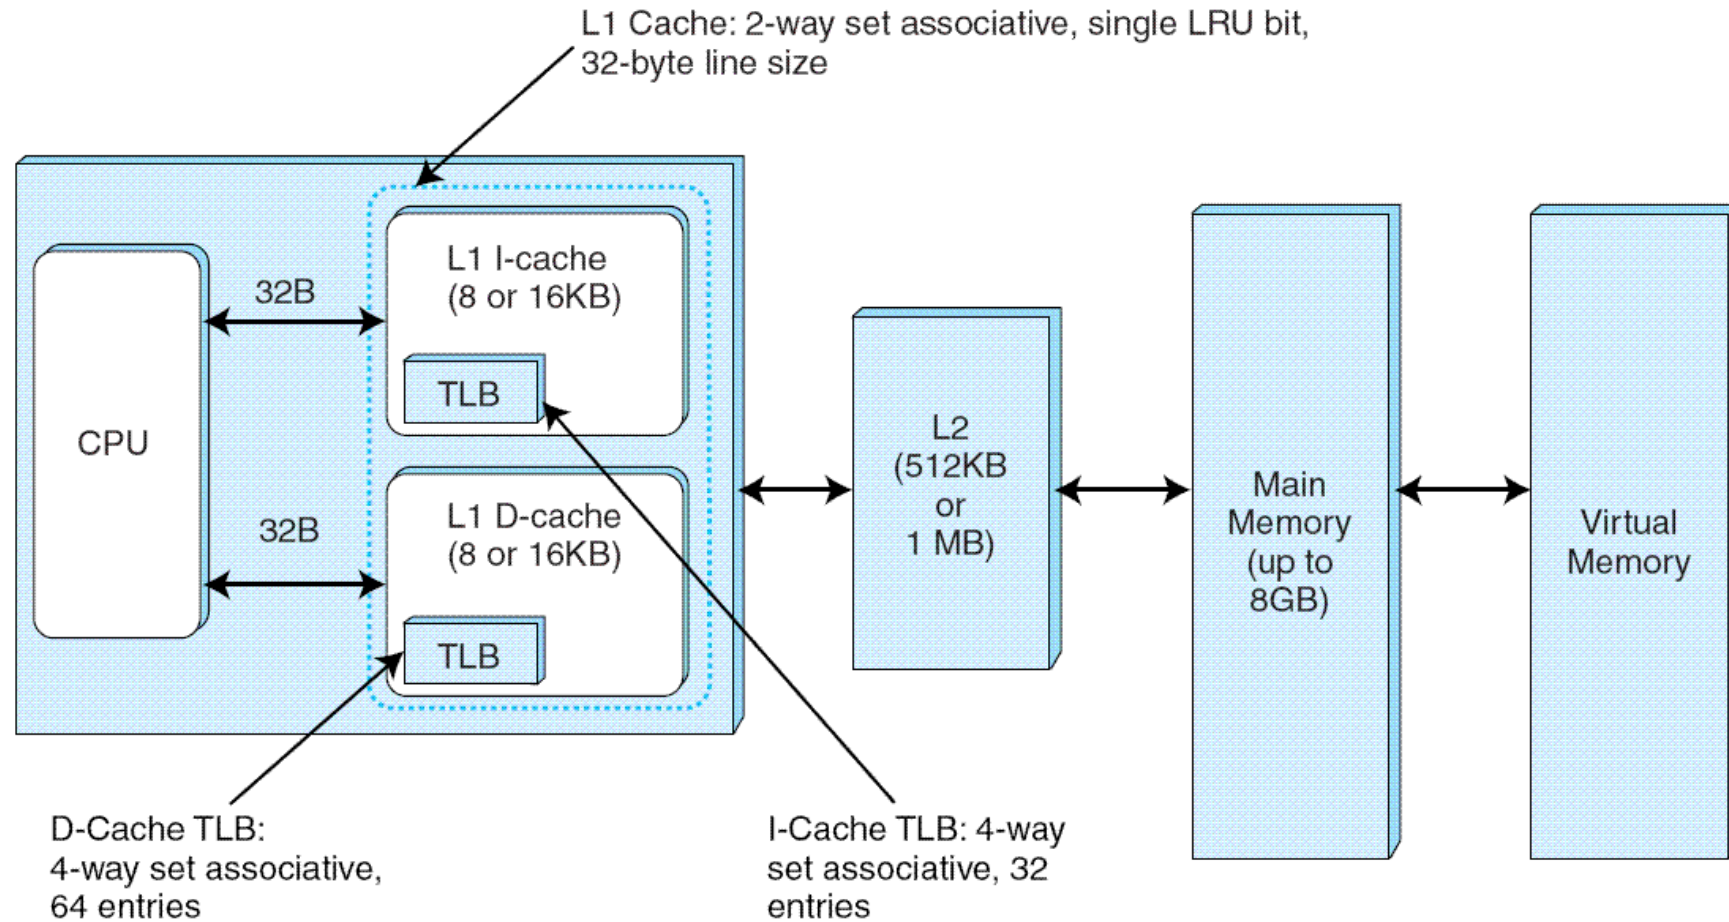
VIRTUAL MEMORY

- Large page tables are cumbersome and slow, but with its uniform memory mapping, page operations are fast. Segmentation allows fast access to the segment table, but segment loading is labor-intensive.
- Paging and segmentation can be combined to take advantage of the best features of both by assigning fixed-size pages within variable-sized segments.
- Each segment has a page table. This means that a memory address will have three fields, one for the segment, another for the page, and a third for the offset.

A REAL-WORLD EXAMPLE

- The Pentium architecture supports both paging and segmentation, and they can be used in various combinations including unpagged unsegmented, segmented unpagged, and unsegmented pagged.
- The processor supports two levels of cache (L1 and L2), both having a block size of 32 bytes.
- The L1 cache is next to the processor, and the L2 cache sits between the processor and memory.
- The L1 cache is in two parts: and instruction cache (I-cache) and a data cache (D-cache).

A REAL-WORLD EXAMPLE



THANK YOU